



UNIVERSITÀ
DEGLI STUDI
DI TRIESTE

Restoration and development of Arm's Java-based LEGv8 ISA simulator

Graduating student - Simone Deiana

Supervisor - Alberto Carini

Restoration and development of Arm's Java-based LEGv8 ISA simulator



Restoration and development of Arm's Java-based LEGv8 ISA simulator






Restoration and development of Arm's Java-based LEGv8 ISA simulator

WHAT IS LEGv8?

Restoration and development of
Arm's Java-based **LEGv8 ISA** simulator

WHAT IS LEGv8?

AN ISA FOR LEARNING COMPUTER ARCHITECTURES



From Computer Organization and Design ARM Edition: The Hardware Software Interface - Patterson, D.A. and Hennessy, J.L.



David A.
Patterson

Peg Skorpinski, CC BY-SA
3.0
<<https://creativecommons.org/licenses/by-sa/3.0/>>, via
Wikimedia Commons



John L.
Hennessy

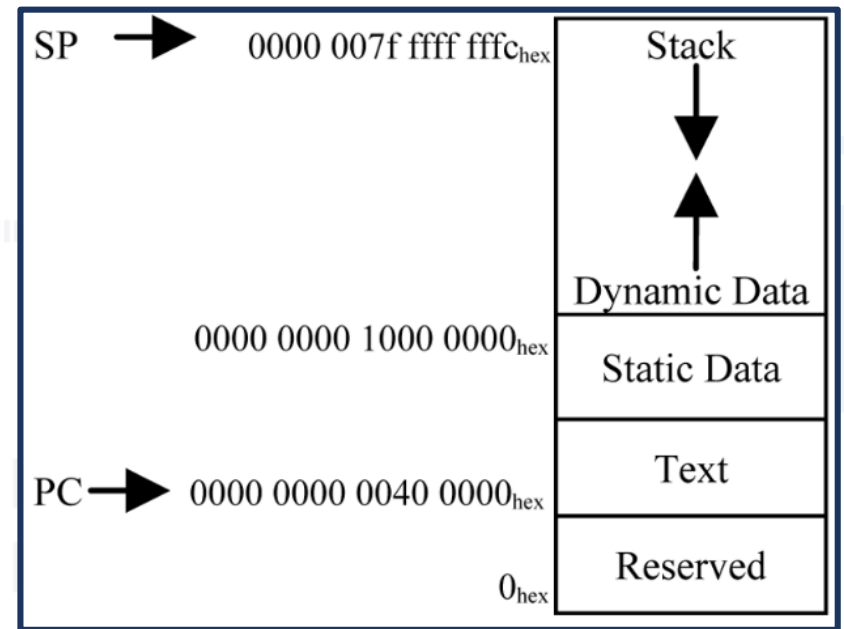
Eric Chan, CC BY 2.0
<<https://creativecommons.org/licenses/by/2.0/>>, via Wikimedia Commons

THE DESIGN PHILOSOPHY

- As simple as it can be...
- ... but with a modern design
- Heavily inspired by ARMv8, almost a "subset"

THE MEMORY

- 64-bit addresses
- Harvard model



THE REGISTERS

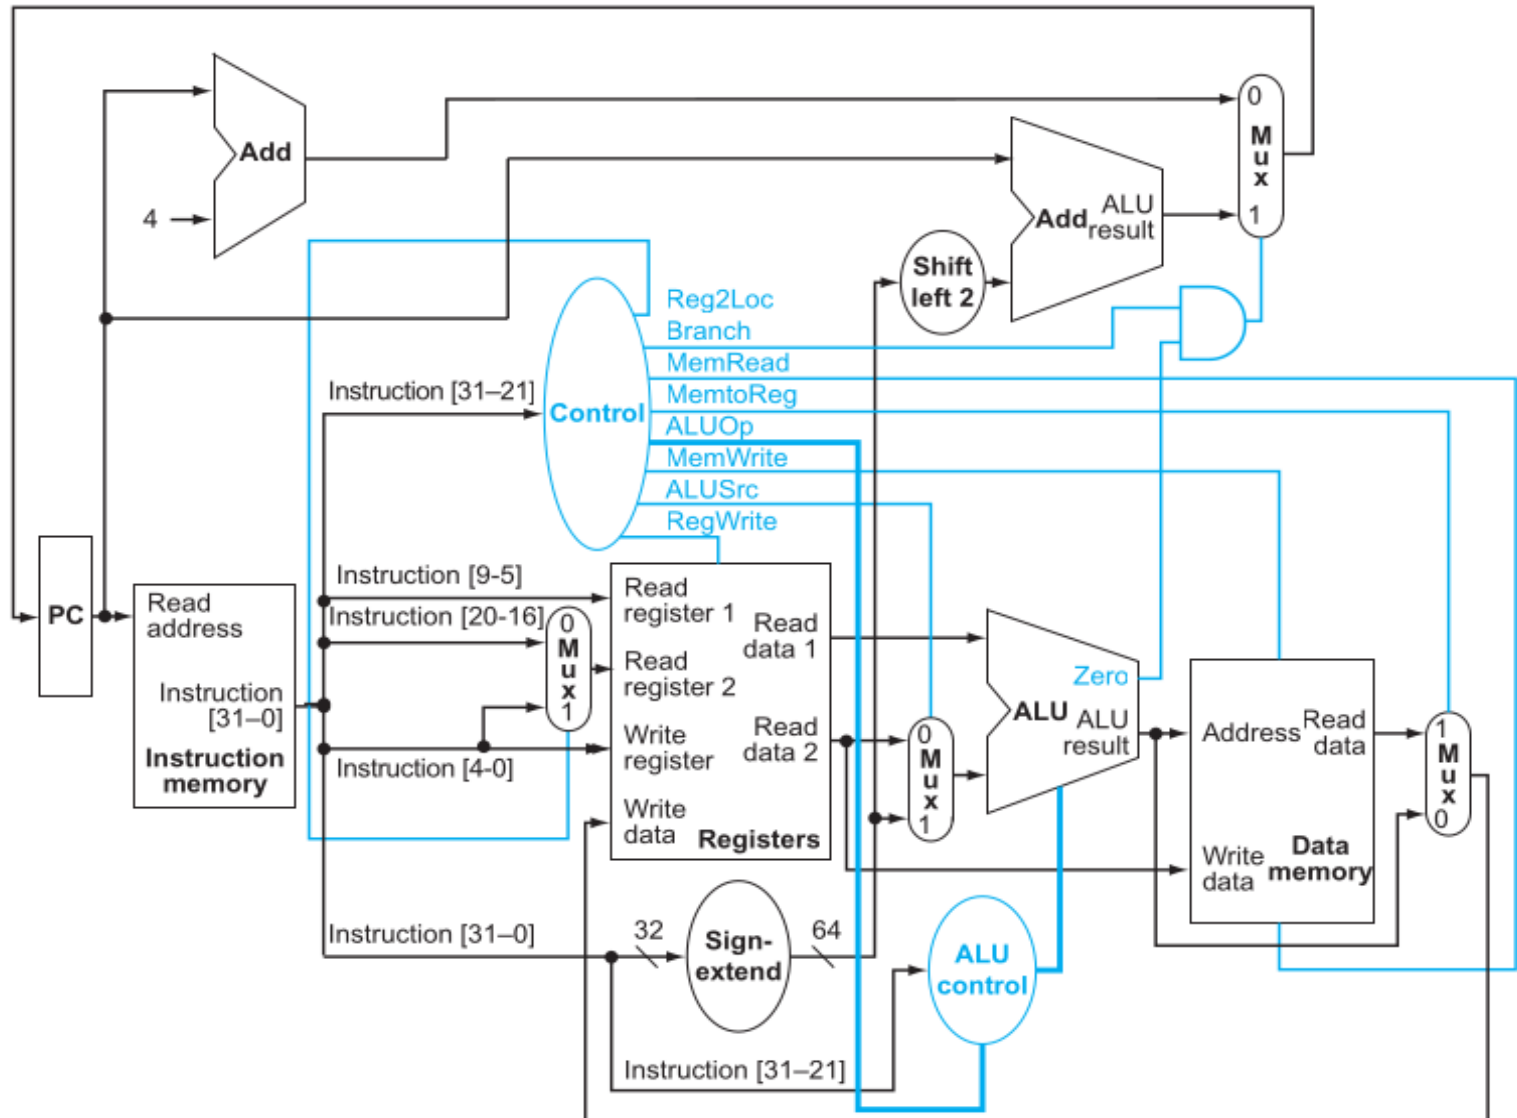
- 32 64-bit "X" integer registers
- 32 64-bit "D" floating-point registers
- 32 32-bit "S" floating-point "registers"

THE INSTRUCTIONS

- 64-bit integer and IEEE-754 floating-point arithmetic
- Designed and optimized for pipelined execution

WHAT IS LEGv8?

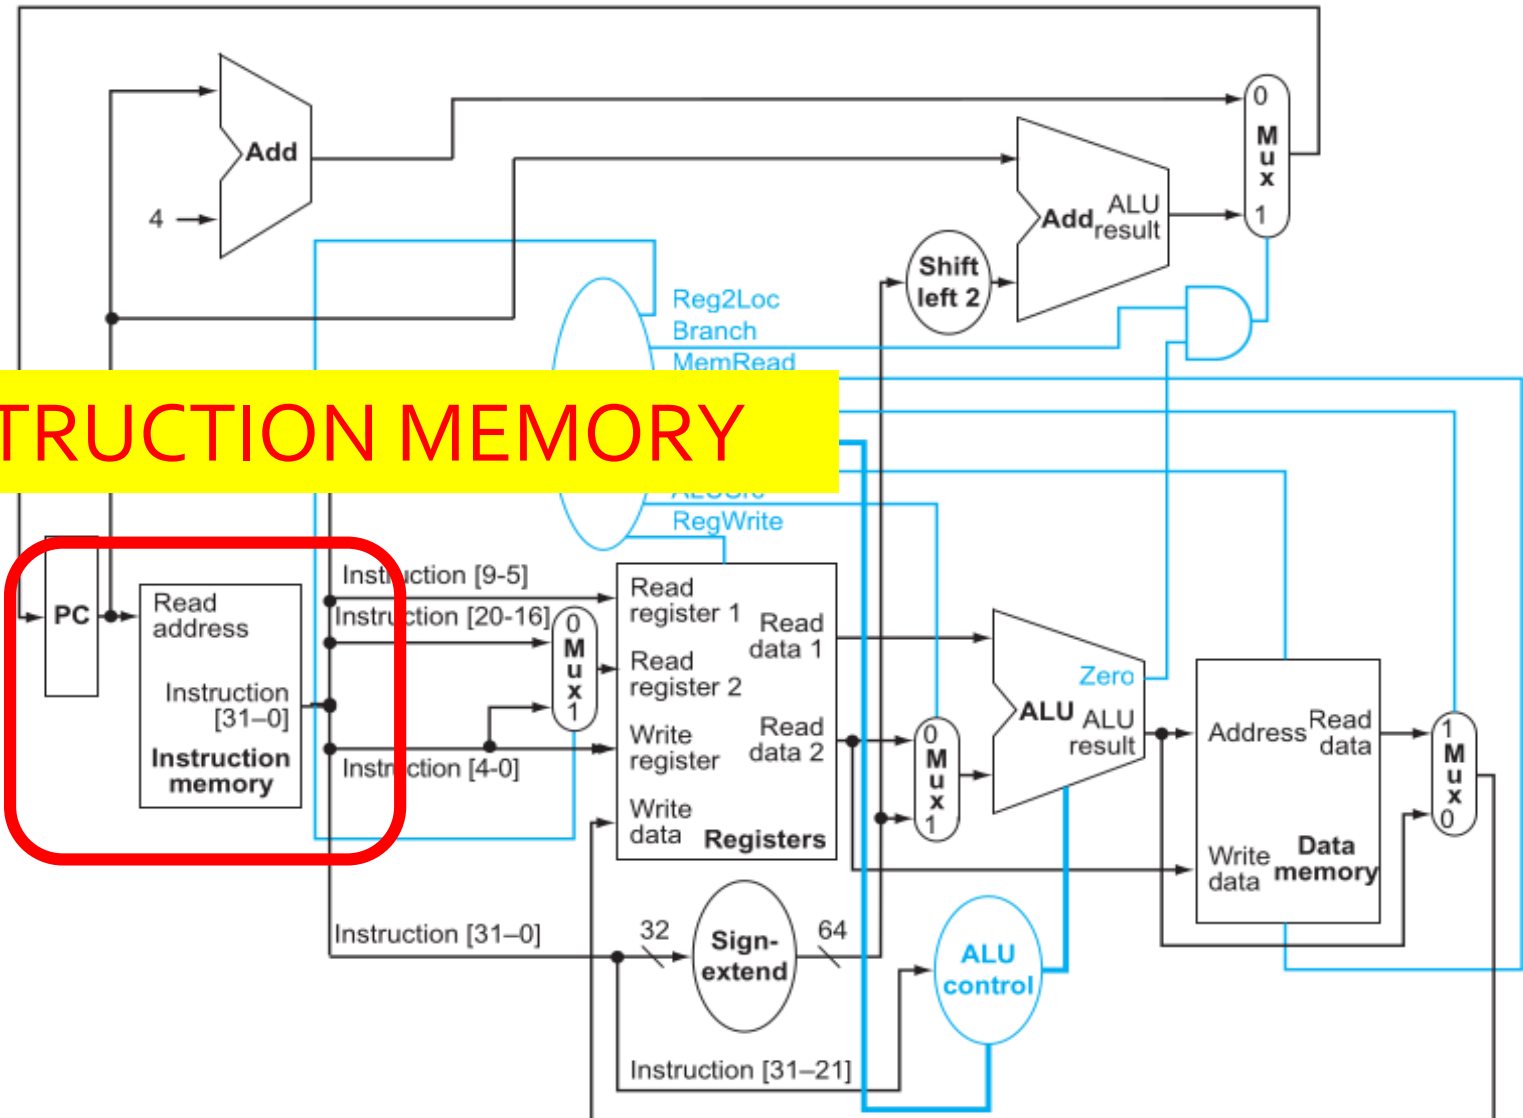
THE DATAPATH



From Computer Organization and Design
ARM Edition: The Hardware Software Interface -
Patterson, D.A. and Hennessy, J.L.

WHAT IS LEGv8?

THE DATAPAH

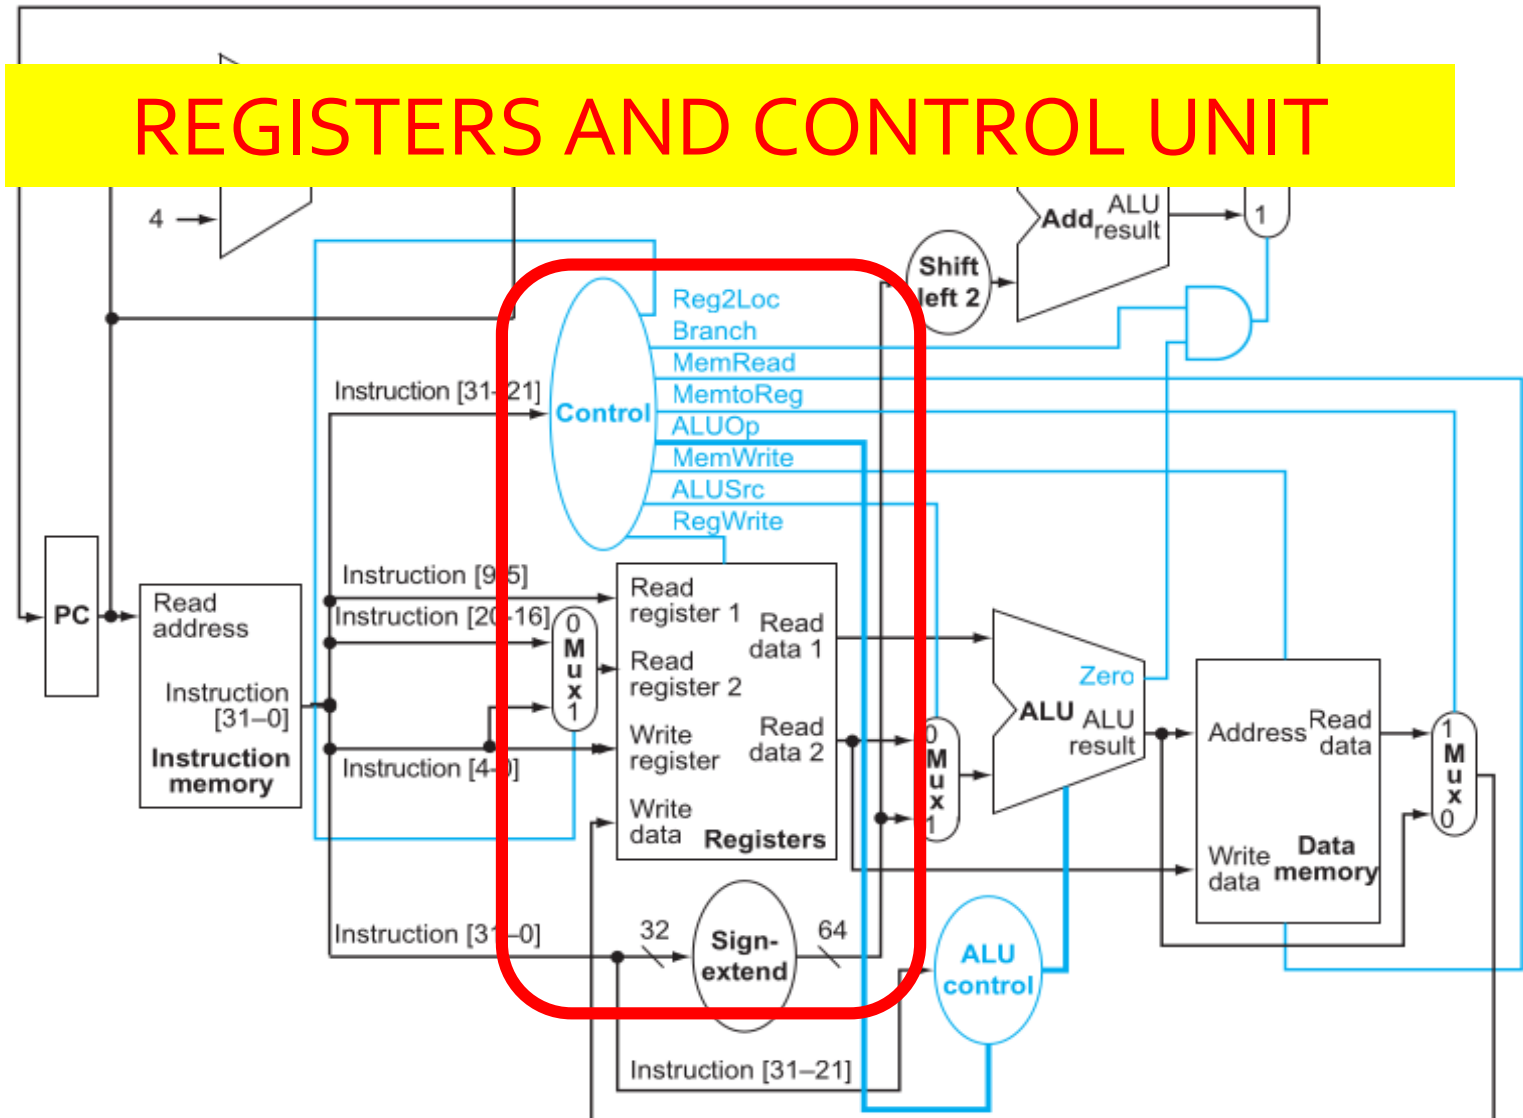


From Computer
Organization and Design
ARM Edition: The Hardware
Software Interface -
Patterson, D.A. and
Hennessy, J.L.

WHAT IS LEGv8?

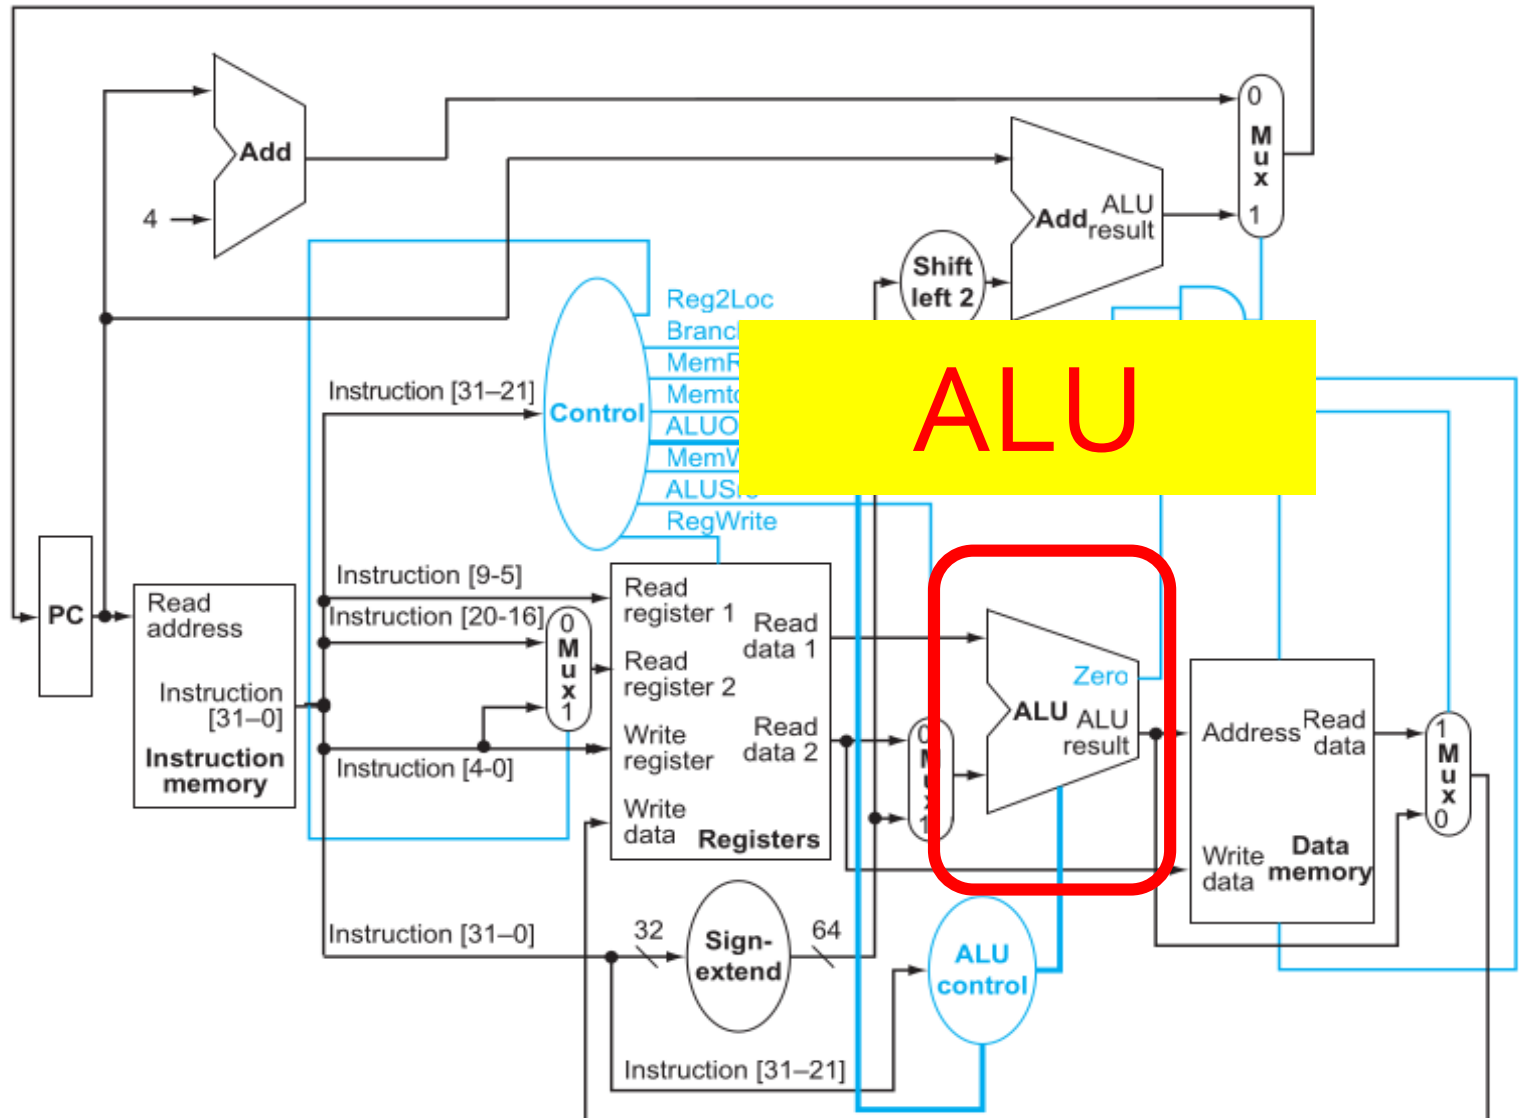
THE DATAPATH

REGISTERS AND CONTROL UNIT



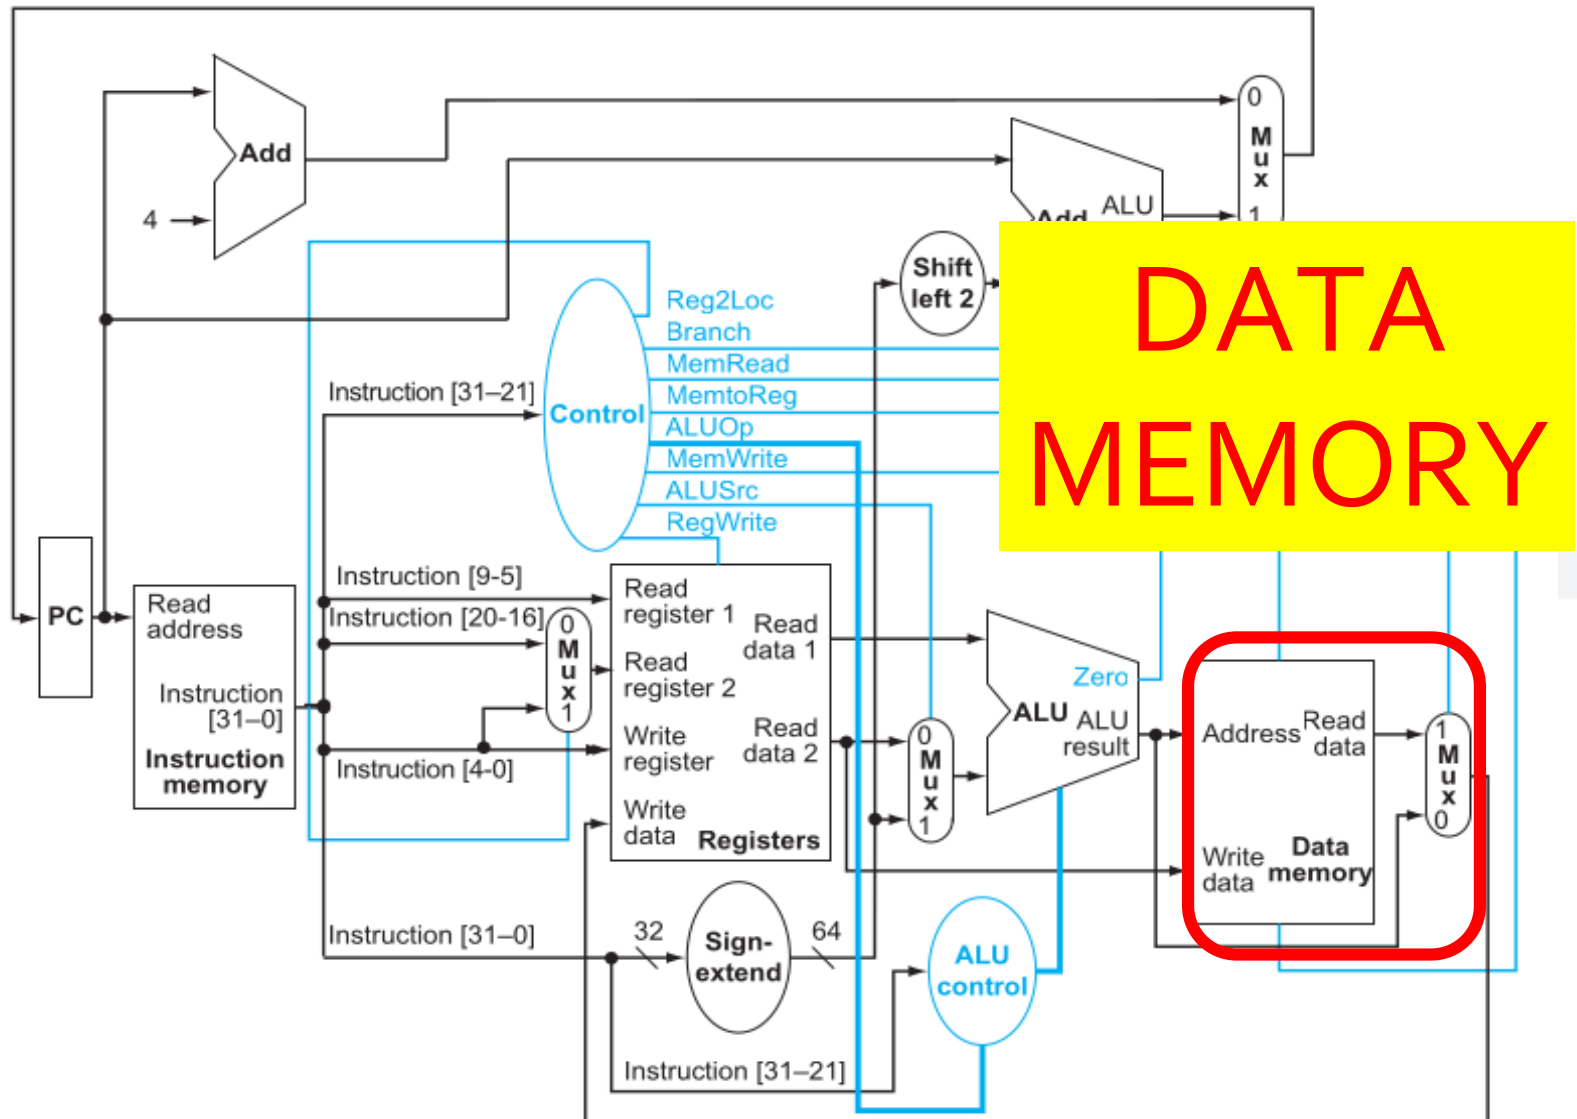
WHAT IS LEGv8?

THE DATAPATH



WHAT IS LEGv8?

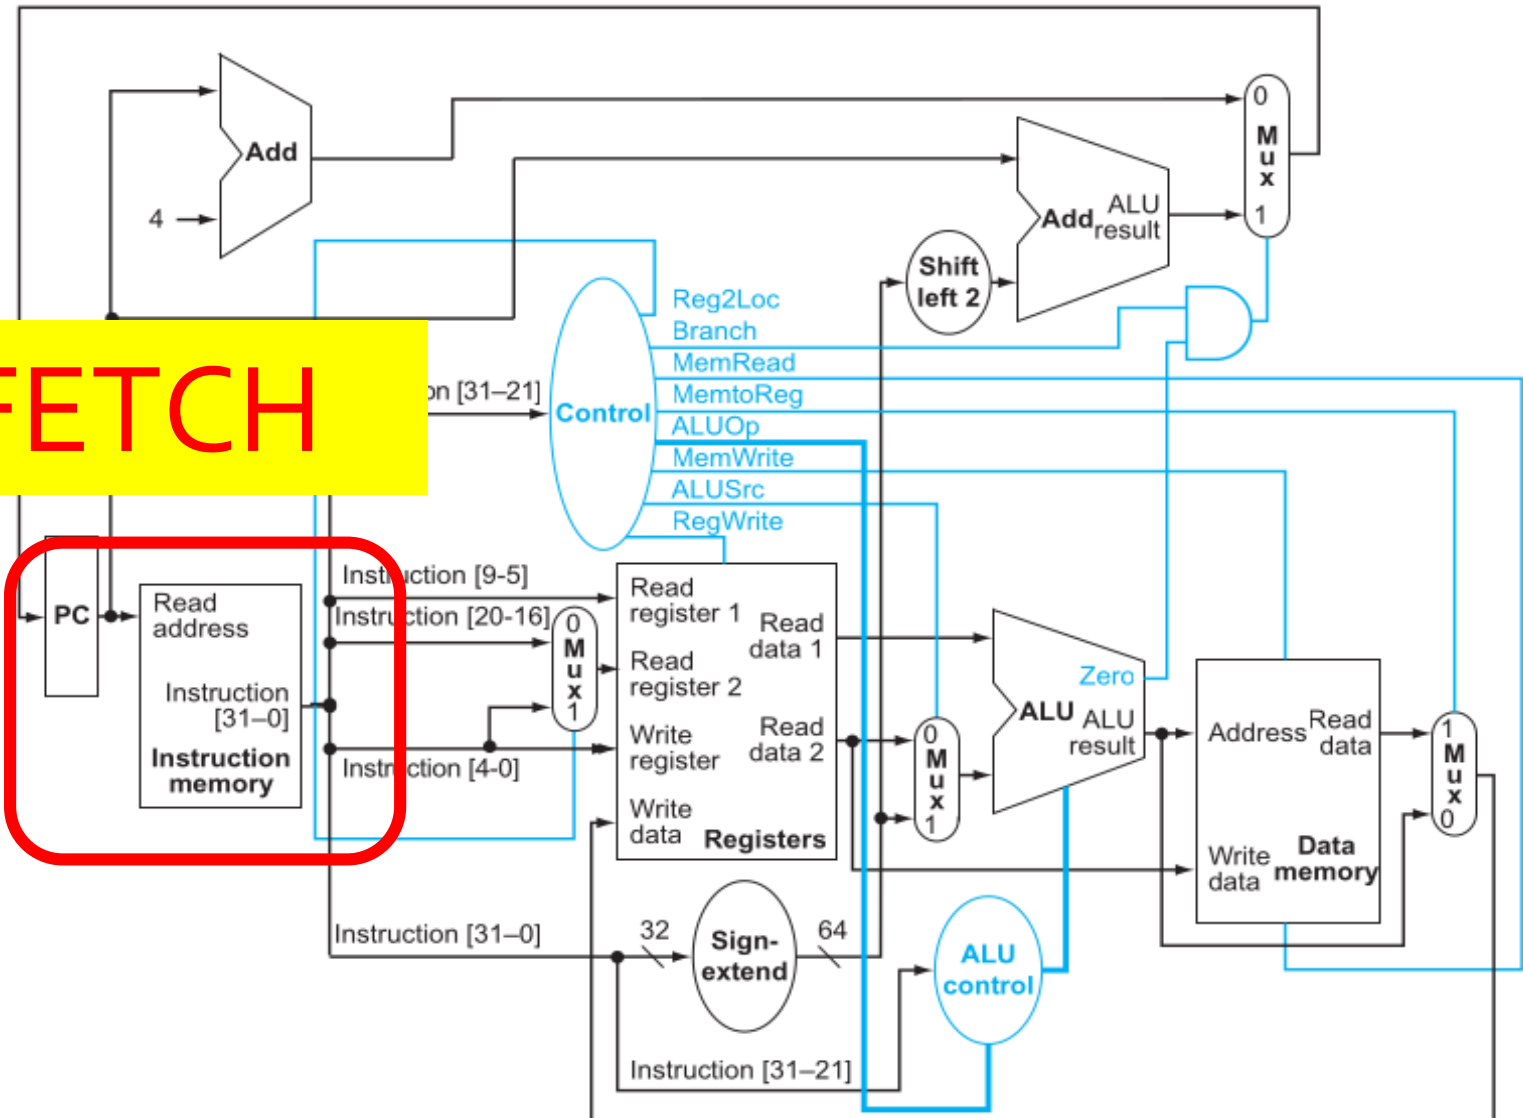
THE DATAPATH



WHAT IS LEGv8?

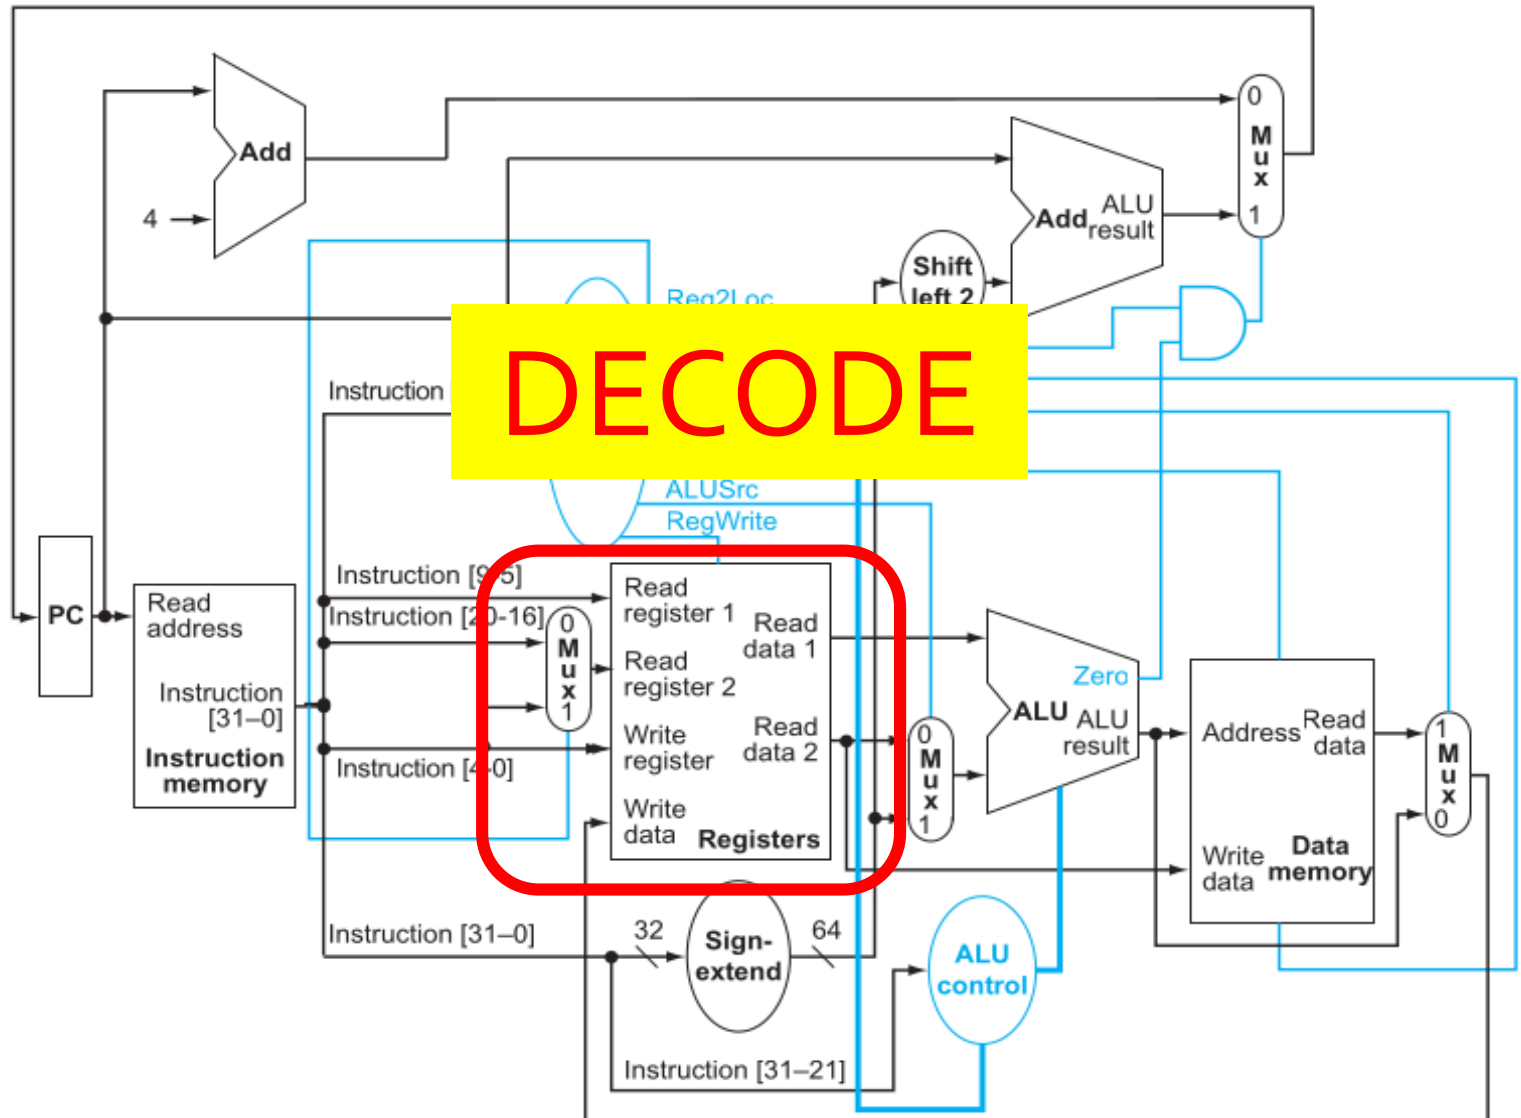
THE STAGES

FETCH



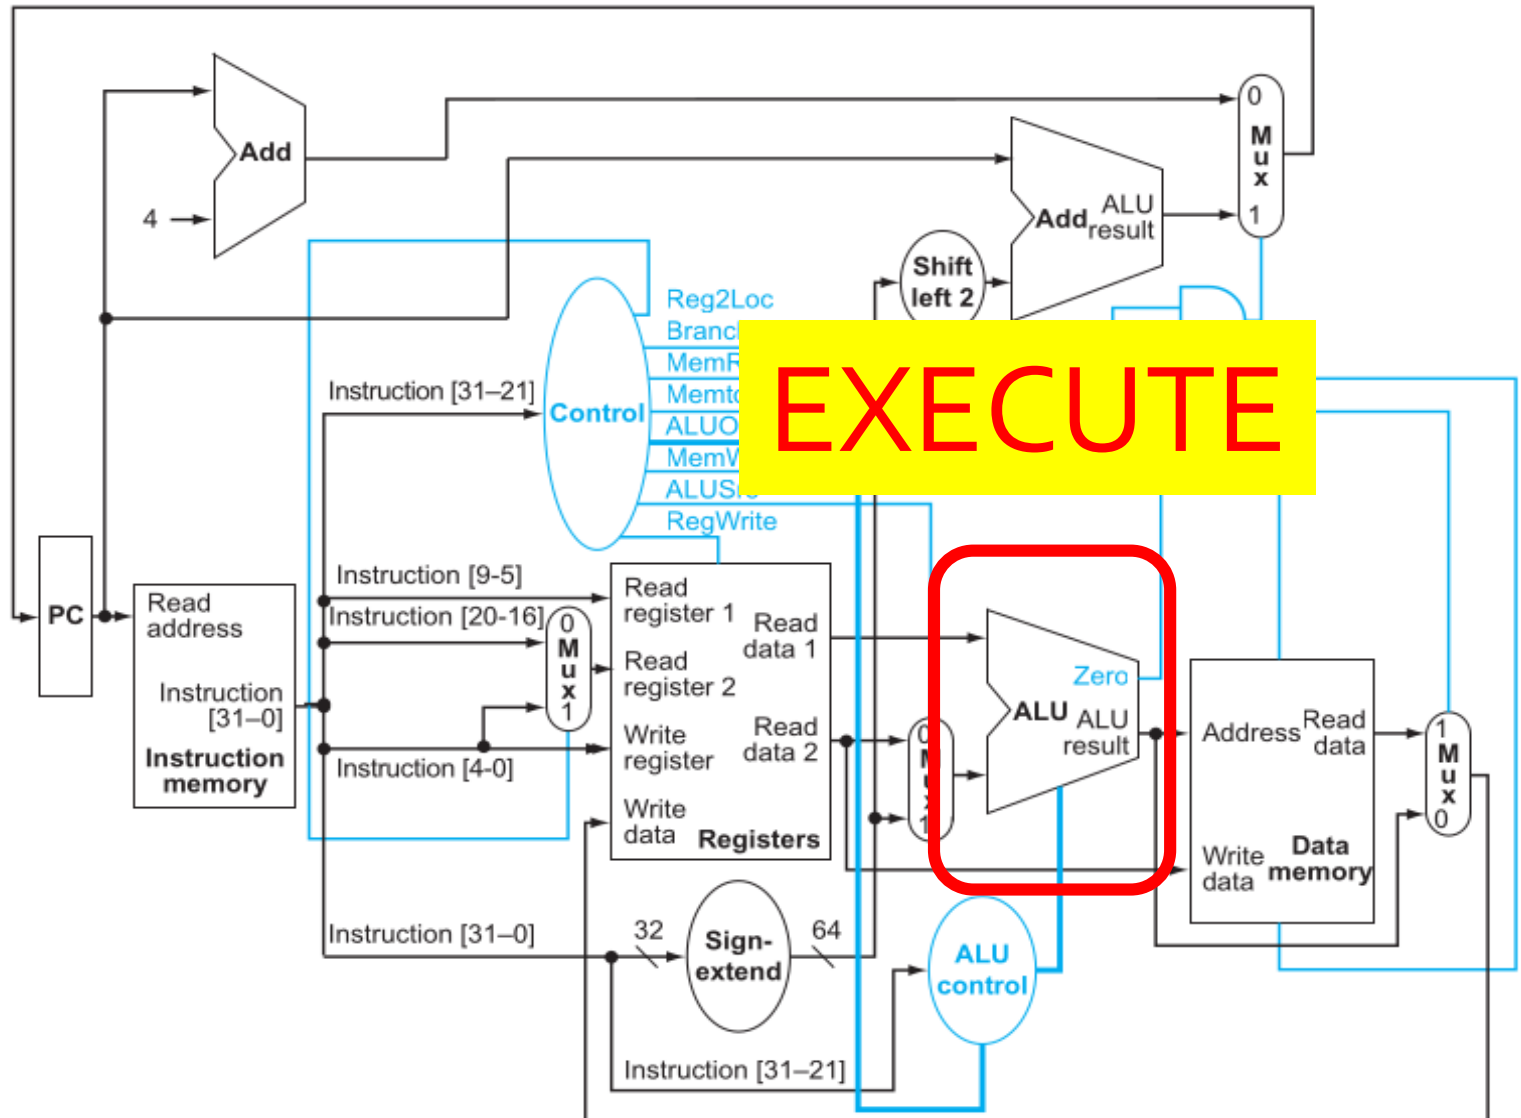
WHAT IS LEGv8?

THE STAGES



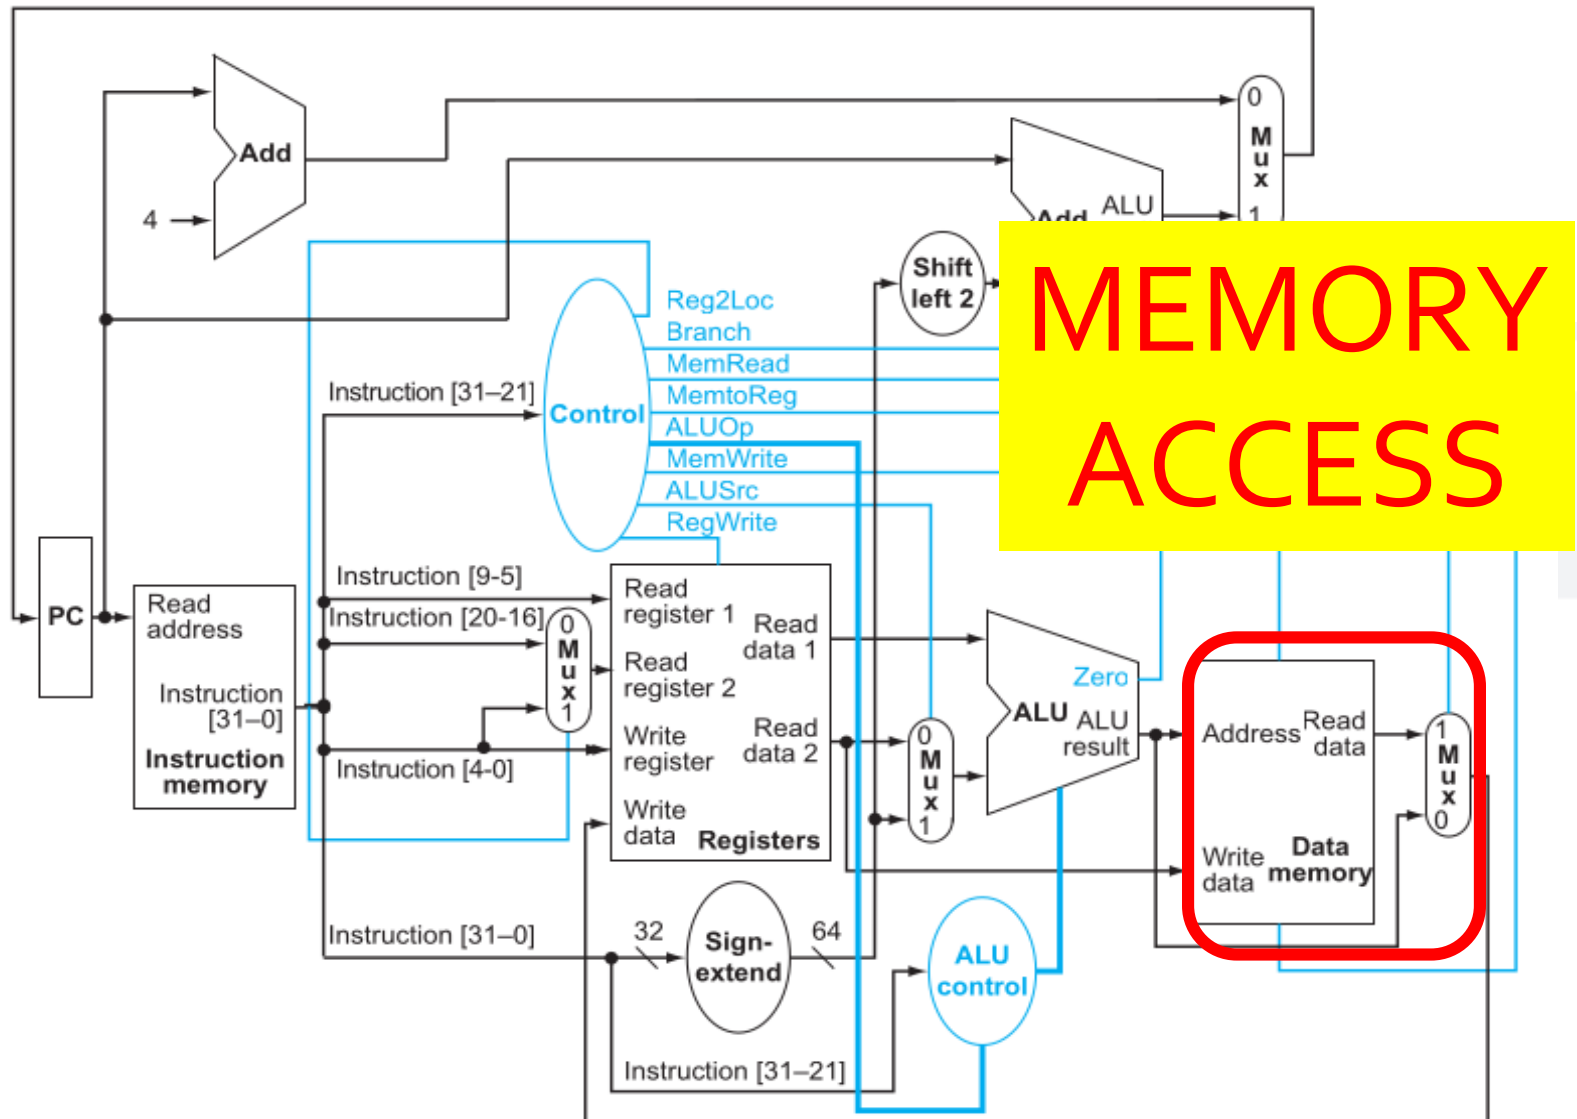
WHAT IS LEGv8?

THE STAGES



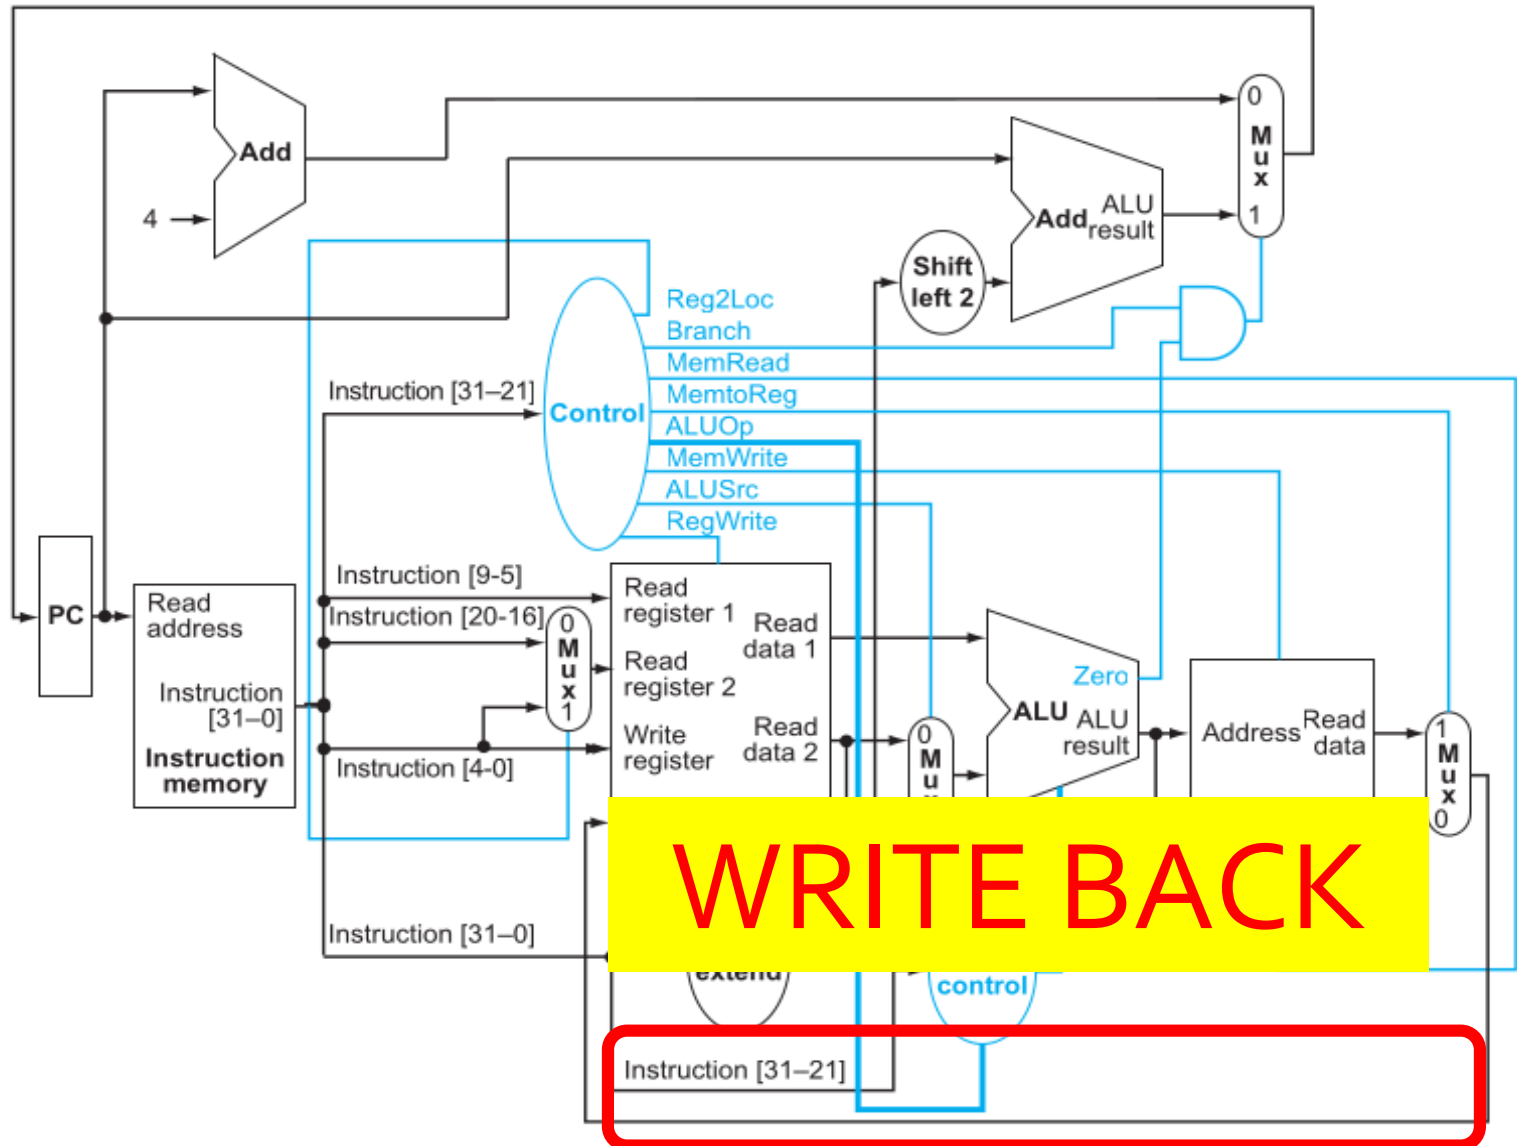
WHAT IS LEGv8?

THE STAGES



WHAT IS LEGv8?

THE STAGES



WHAT SIMULATOR, AND WHY?

Restoration and development of
Arm's Java-based LEGv8 ISA simulator

WHAT SIMULATOR, AND WHY?

NO HARDWARE FOR LEGv8 => NEED A SIMULATOR



Michael H. („Laserlicht“) / Wikimedia Commons

WHAT SIMULATOR, AND WHY?

BUT WHICH ONE?



LEGv8

197 results

THE PROBLEM:

NO SOFTWARE CAN YET
SIMULATE THE ENTIRE
LEGv8 ISA!

THE SOLUTION:

- Write one from scratch
- OR (BETTER)
- Improve one that already exists

ARM HAS OFFICIALLY MADE A LEGv8 SIMULATOR

GOOD!

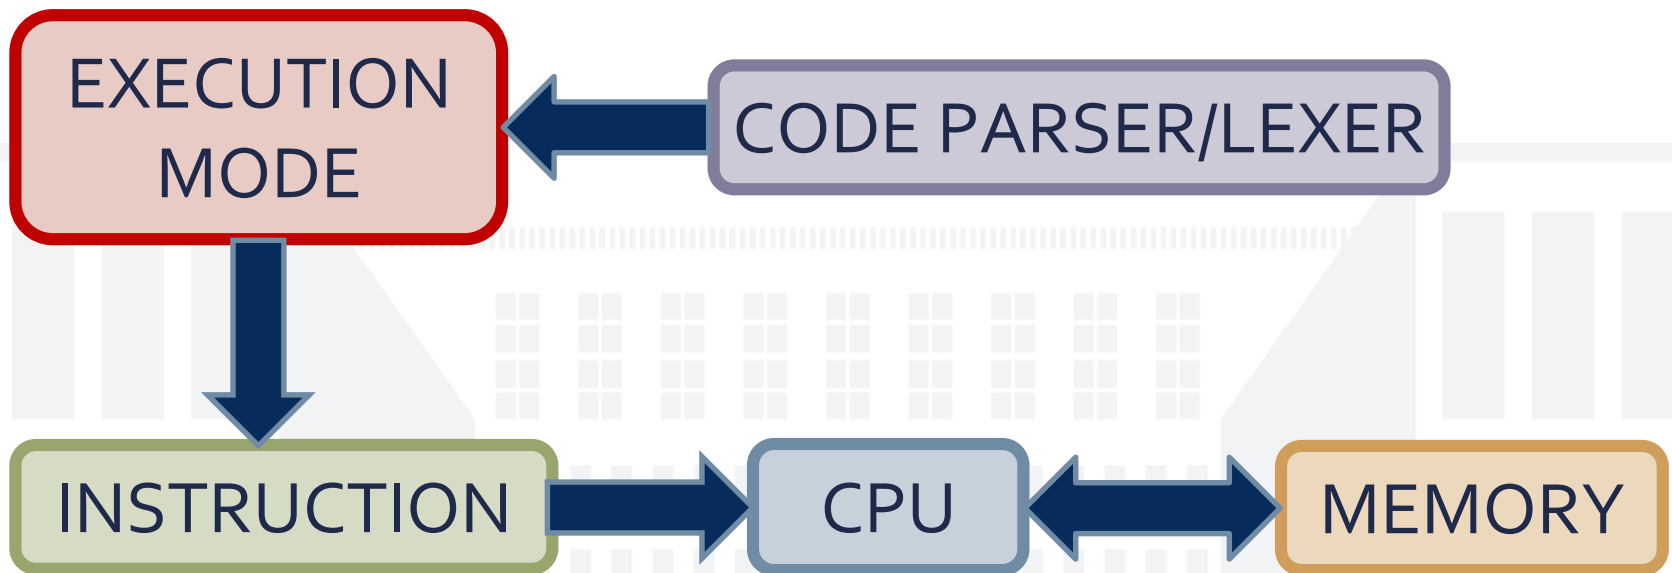
IT'S INCOMPLETE
AND BROKEN

BAD!

WHAT STANDS OUT:

- Written in Java (high level, extensible)
- Distributed as a web application
- Nice, functional UI
- Closely follows the textbook

STRUCTURE OF THE SIMULATOR



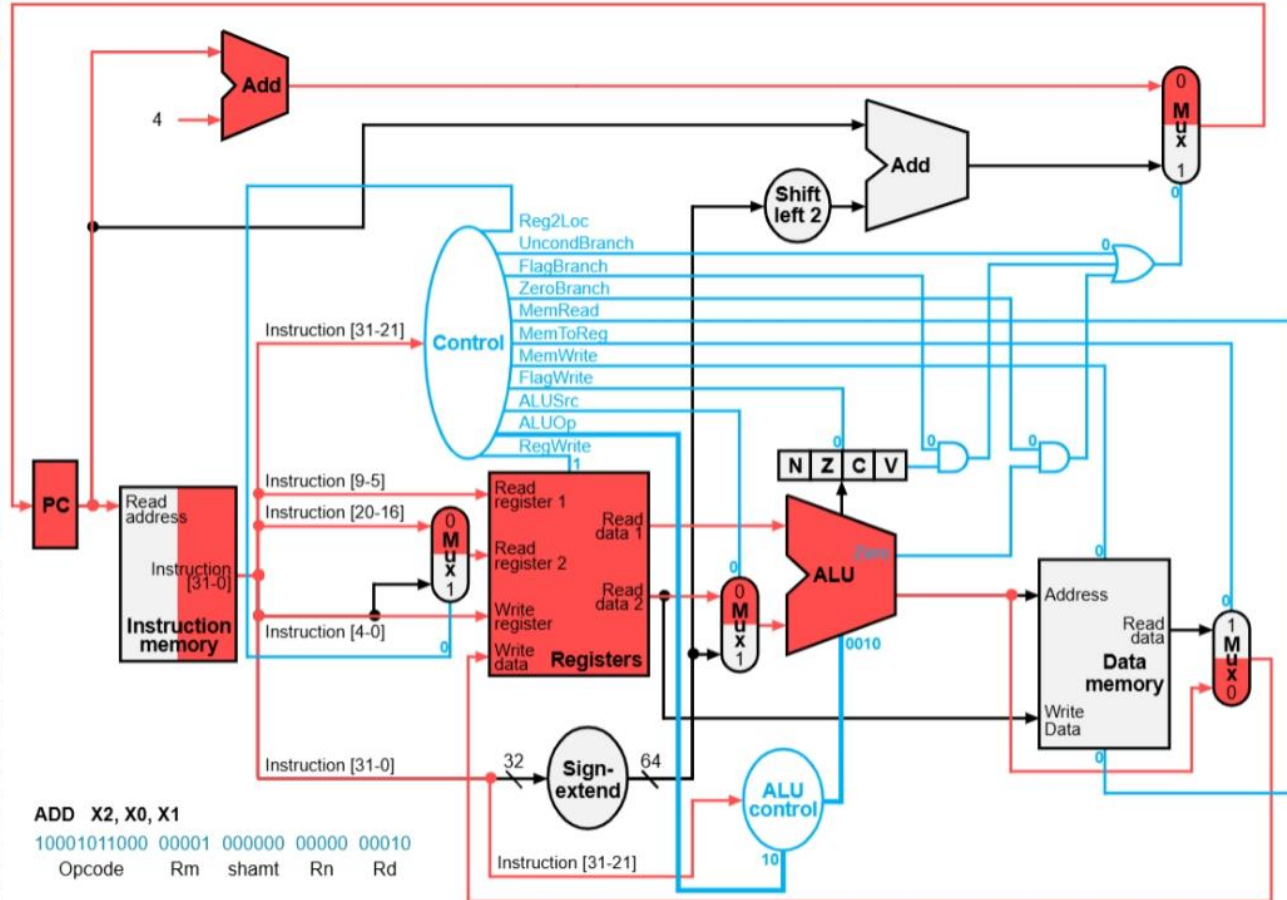
WHAT SIMULATOR, AND WHY?

LEGv8 Simulator

Execution Mode: **Single Cycle** Assemble Execute Instruction Help

```
1 ADDI X0, XZR, #53
2 ADDI X1, XZR, #75
3 ADD X2, X0, X1
4 CMP X2, X0
5 B.LE if
6 if: SUBIS X3, X2, #9
7 MOV X4, X3
8
```

PC	0x40000c	Hex	Dec	Z	0	N	0	C	0	V	0
X0	0x35	Hex	Dec	X16	0x0	Hex	Dec	X16	0x0	Hex	Dec
X1	0x4b	Hex	Dec	X17	0x0	Hex	Dec	X17	0x0	Hex	Dec
X2	0x80	Hex	Dec	X18	0x0	Hex	Dec	X18	0x0	Hex	Dec
X3	0x0	Hex	Dec	X19	0x0	Hex	Dec	X19	0x0	Hex	Dec
X4	0x0	Hex	Dec	X20	0x0	Hex	Dec	X20	0x0	Hex	Dec
X5	0x0	Hex	Dec	X21	0x0	Hex	Dec	X21	0x0	Hex	Dec
X6	0x0	Hex	Dec	X22	0x0	Hex	Dec	X22	0x0	Hex	Dec
X7	0x0	Hex	Dec	X23	0x0	Hex	Dec	X23	0x0	Hex	Dec
X8	0x0	Hex	Dec	X24	0x0	Hex	Dec	X24	0x0	Hex	Dec
X9	0x0	Hex	Dec	X25	0x0	Hex	Dec	X25	0x0	Hex	Dec
X10	0x0	Hex	Dec	X26	0x0	Hex	Dec	X26	0x0	Hex	Dec
X11	0x0	Hex	Dec	X27	0x0	Hex	Dec	X27	0x0	Hex	Dec
X12	0x0	Hex	Dec	SP	0x7fffffc	Hex	Dec	SP	0x7fffffc	Hex	Dec
X13	0x0	Hex	Dec	FP	0x0	Hex	Dec	FP	0x0	Hex	Dec
X14	0x0	Hex	Dec	LR	0x0	Hex	Dec	LR	0x0	Hex	Dec
X15	0x0	Hex	Dec	XZR	0x0	Hex	Dec	XZR	0x0	Hex	Dec



FIXING AND RESTORING THE SIMULATOR

Restoration and development of
Arm's Java-based LEGv8 ISA simulator

COMPARISONS DON'T WORK!

- No "if-else" conditionals
- No "switch-case" conditionals
- No "while" loops
- No "for" loops

THE COMPARISON BUG

```
1      ADDI    X0, XZR, #1
2      ADDI    X1, XZR, #2
3      CMP     X0, X1
4      B.LE    should_jump
5      ADD     X2, X0, X1
6  should_jump:  ADDI    X2, XZR, #16
7
```

X0 = 1, X1 = 2


PC	0x400008	Hex	Dec	Z	0	N	0	C	0	V	0
X0	0x1	Hex	Dec	X16			0x0			Hex	Dec
X1	0x2	Hex	Dec	X17			0x0			Hex	Dec
X2	0x0	Hex	Dec	X18			0x0			Hex	Dec

THE COMPARISON BUG

```
1 |      ADDI      X0, XZR, #1
2 |      ADDI      X1, XZR, #2
3 | CMP           X0, X1
4 |      B.LE      should_jump
5 |      ADD       X2, X0, X1
6 | should_jump:   ADDI      X2, XZR, #16
7 |
```

COMPARE X0 WITH X1

PC	0x40000c	Hex	Dec	Z	0	N	1	C	0	V	1
X0	0x1	Hex	Dec	X16			0x0			Hex	Dec
X1	0x2	Hex	Dec	X17			0x0			Hex	Dec
X2	0x0	Hex	Dec	X18			0x0			Hex	Dec



THE COMPARISON BUG

```
1 |      ADDI      X0, XZR, #1
2 |      ADDI      X1, XZR, #2
3 |      CMP       X0, X1
4 |      B.LE      should_jump
5 |      ADD       X2, X0, X1
6 |  should_jump:  ADDI      X2, XZR, #16
7 |
```

1 <= 2 ?

OF COURSE, LET'S JUMP!

PC	0x400010	Hex	Dec	Z	0	N	1	C	0	V	1
X0	0x1	Hex	Dec	X16			0x0	Hex	Dec		
X1	0x2	Hex	Dec	X17			0x0	Hex	Dec		
X2	0x0	Hex	Dec	X18			0x0	Hex	Dec		

THE COMPARISON BUG

```
1 |      ADDI      X0, XZR, #1
2 |      ADDI      X1, XZR, #2
3 |      CMP       X0, X1
4 |      B.LE      should_jump
5 |      ADD       X2, X0, X1
6 | should_jump:   ADDI      X2, XZR, #16
7 |
```

...OR NOT

PC	0x400014	Hex	Dec	Z	0	N	1	C	0	V	1
X0	0x1	Hex	Dec	X16			0x0	Hex	Dec		
X1	0x2	Hex	Dec	X17			0x0	Hex	Dec		
X2	0x3	Hex	Dec	X18			0x0	Hex	Dec		

THE FLAGS ARE SET WRONG!

HOW THEY SHOULD BE

\leq	B.LE	$\sim(Z=0 \ \& \ N=V)$
--------	------	------------------------

HOW THEY ARE

Z	0	N	1	C	0	V	1
---	---	---	---	---	---	---	---

$\neg(Z=0 \ \wedge \ N=V) = \neg(\text{TRUE} \ \wedge \ \text{TRUE}) = \text{FALSE}$

BRANCH AND LINKS DON'T WORK!

~~*void subroutine(arg1, ...)*~~
~~*float function(arg1, ...)*~~

CAN'T REUSE CODE

THE BRANCH AND LINK BUG

```
1      ADDI      X0, XZR, #1
2      BL        subroutine
3      B         exit
4
5  subroutine:
6      ADDI      X0, X0, #16
7      BR        LR
8  exit:
9
```

PROGRAM COUNTER: 0x0

RETURN ADDRESS: 0x0

THE BRANCH AND LINK BUG

```
1      ADDI      X0, XZR, #1
2      BL        subroutine
3      B         exit
4
5  subroutine:
6      ADDI      X0, X0, #16
7      BR        LR
8  exit:
9
```

PROGRAM COUNTER: 0x4

RETURN ADDRESS: 0x0

THE BRANCH AND LINK BUG

```
1      ADDI    X0, XZR, #1
2      BL      subroutine
3      B       exit
4
5  subroutine:
6      ADDI    X0, X0, #16
7      BR      LR
8  exit:
9
```

PROGRAM COUNTER: 0xC

RETURN ADDRESS: 0xC

THE BRANCH AND LINK BUG

```
1      ADDI      X0, XZR, #1
2      BL        subroutine
3      B         exit
4
5  subroutine:
6      ADDI      X0, X0, #16
7      BR        LR
8  exit:
9
```

PROGRAM COUNTER: 0x10

RETURN ADDRESS: 0xC

THE BRANCH AND LINK BUG

```
1      ADDI      X0, XZR, #1
2      BL        subroutine
3      B         exit
4
5  subroutine:
6      ADDI      X0, X0, #16
7      BR        LR
8  exit:
9
```

IT DOESN'T GO BACK!

IT SHOULD GO HERE!

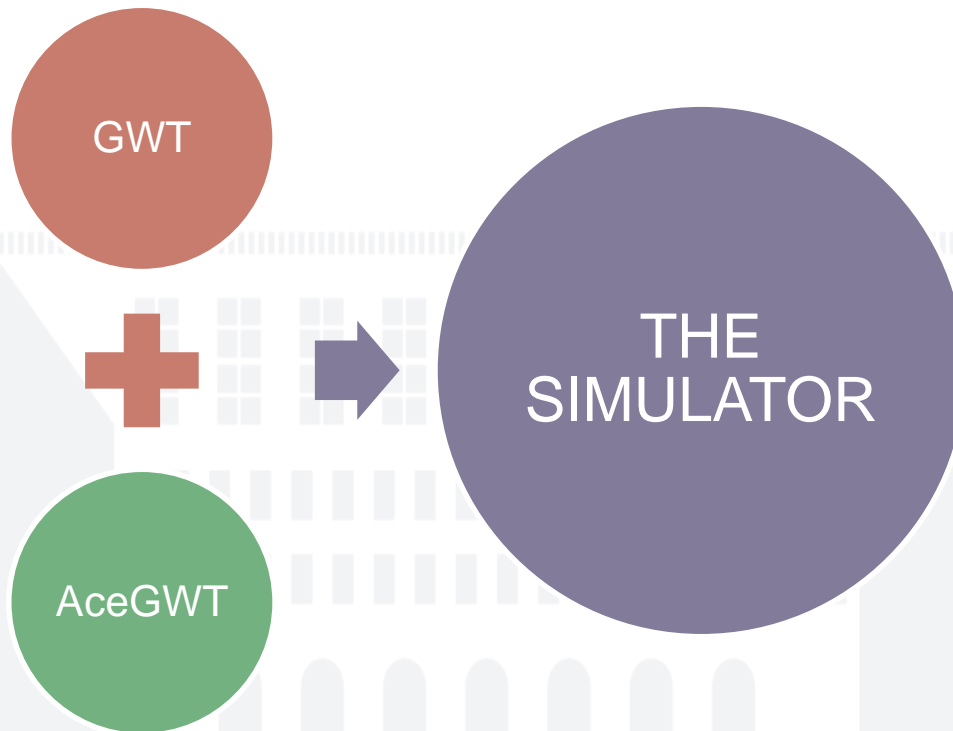
1	ADDI	X0, XZR, #1	
2	BL	subroutine	
3	B	exit	ADDRESS: 0x8
4			
5	subroutine:		
6	ADDI	X0, X0, #16	
7	BR	LR	
8	exit:		
9			

PROGRAM COUNTER: 0xC

RETURN ADDRESS: 0xC

ALL FIXED, BUT...
NOBODY KNOWS HOW IT WORKS!

THE PROJECT'S DEPENDENCIES



GWT

- Framework (formerly) from Google
- Generates web applications (client-server, client only) from Java
- Emulates Java's JVM with JavaScript

GWT

- Old, outdated, barely supported
- Convoluted custom build tools
- Limited emulation of JVM
- Basically needs Eclipse plug-in for real development

AceGWT

- Provides GWT bindings for the Ace editor
- Can be used like normal GWT component
- Also old, outdated, and unsupported

WORKING IT OUT:

- Need old version of Eclipse, and Eclipse GWT plug-in
- Reverse engineer the dependencies and where they are needed
- Configure the project to stop failing

FILLING THE GAPS

Restoration and development of
Arm's Java-based LEGv8 ISA simulator

WHAT IS THE SIMULATOR MISSING?

- Incomplete integer arithmetic
- No IEEE-754 arithmetic and data instructions
- No visualization for the stack memory

THE MISSING INTEGER-BASED INSTRUCTIONS

- **MUL** — LOWER 64 BITS OF THE MULTIPLICATION
- **SMULH** — HIGHER 64 BITS OF THE SIGNED MULTIPLICATION
- **UMULH** — HIGHER 64 BITS OF THE UNSIGNED MULTIPLICATION
- **SDIV** — SIGNED DIVISION
- **UDIV** — UNSIGNED DIVISION
- **LDA** — LOAD ADDRESS OF A LABEL IN A REGISTER

UMULH: A CASE STUDY

- Takes two 64-bit unsigned integer values
- Extends them to 128 bits unsigned
- Performs 128-bit product
- Saves higher 64 bits to destination

PROBLEM 1

- Java does not have primitive 128-bit integer types
- Product of 64-bit integers truncated
- The BigInteger library exists
- GWT 2.7 doesn't emulate it (2.8 does)

PROBLEM 2

- Primitive integers are signed
- BigInteger also signed
- Bitmask converts 64-bit unsigned integers to 65-bit signed, perform signed multiplication, take the higher bits

CAN'T SEE THE STACK

- Fundamental for testing and debugging complex programs (now we can write them)
- Useful to understanding LEGv8 and stack management
- Visible in most simulators

THE INTEGER REGISTERS VIEW

X0	0x0	Hex	Dec	X16	0x0	Hex	Dec
X1	0x0	Hex	Dec	X17	0x0	Hex	Dec
X2	0x0	Hex	Dec	X18	0x0	Hex	Dec
X3	0x0	Hex	Dec	X19	0x0	Hex	Dec
X4	0x0	Hex	Dec	X20	0x0	Hex	Dec
X5	0x0	Hex	Dec	X21	0x0	Hex	Dec
X6	0x0	Hex	Dec	X22	0x0	Hex	Dec
X7	0x0	Hex	Dec	X23	0x0	Hex	Dec
X8	0x0	Hex	Dec	X24	0x0	Hex	Dec
X9	0x0	Hex	Dec	X25	0x0	Hex	Dec
X10	0x0	Hex	Dec	X26	0x0	Hex	Dec
X11	0x0	Hex	Dec	X27	0x0	Hex	Dec
X12	0x0	Hex	Dec	SP	0x7fffffff	Hex	Dec
X13	0x0	Hex	Dec	FP	0x0	Hex	Dec
X14	0x0	Hex	Dec	LR	0x0	Hex	Dec
X15	0x0	Hex	Dec	XZR	0x0	Hex	Dec

THE NEW STACK VIEW

0x8000000000:	0x0	Hex	0x7fffffffff80:	0x0	Hex
0x7fffffffff8:	0x0	Hex	0x7fffffffff78:	0x0	Hex
0x7fffffffff0:	0x0	Hex	0x7fffffffff70:	0x0	Hex
0x7fffffffffe8:	0x0	Hex	0x7fffffffff68:	0x0	Hex
0x7fffffffffe0:	0x0	Hex	0x7fffffffff60:	0x0	Hex
0x7ffffffffffd8:	0x0	Hex	0x7ffffffffff58:	0x0	Hex
0x7ffffffffffd0:	0x0	Hex	0x7ffffffffff50:	0x0	Hex
0x7ffffffffffc8:	0x0	Hex	0x7ffffffffff48:	0x0	Hex
0x7ffffffffffc0:	0x0	Hex	0x7ffffffffff40:	0x0	Hex
0x7ffffffffffb8:	0x0	Hex	0x7ffffffffff38:	0x0	Hex
0x7ffffffffffb0:	0x0	Hex	0x7ffffffffff30:	0x0	Hex
0x7ffffffffffa8:	0x0	Hex	0x7ffffffffff28:	0x0	Hex
0x7ffffffffffa0:	0x0	Hex	0x7ffffffffff20:	0x0	Hex
0x7ffffffffff98:	0x0	Hex	0x7ffffffffff18:	0x0	Hex
0x7ffffffffff90:	0x0	Hex	0x7ffffffffff10:	0x0	Hex
0x7ffffffffff88:	0x0	Hex	0x7ffffffffff08:	0x0	Hex

ADDING FLOATING-POINT SUPPORT

- **FADDS, FADDD** — ADD TWO IEEE-754 VALUES
- **FSUBS, FSUBD** - SUBTRACT TWO IEEE-754 VALUES
- **FMULS, FMULD** - MULTIPLY TWO IEEE-754 VALUES
- **FDIVS, FDIVD** — DIVIDE TWO IEEE-754 VALUES
- **LDURS, LDURD** — LOAD IEEE-754 VALUE FROM MEMORY
- **STURS, STURD** - STORE IEEE-754 VALUE TO MEMORY
- **FCMPS, FCMPPD** — COMPARE TWO IEEE-754 VALUES

ARITHMETICAL INSTRUCTIONS (FADDD, FDIVS, ...)

- Native Java support for IEEE-754 with float and double types
- Native Java support for IEEE-754 arithmetical operations
- Straight forward implementation

MEMORY ACCESS INSTRUCTIONS (LDURS, STURD, ...)

- Simulator uses long values to store bits in memory
- Use existing *longs* and *ints* as raw bits
- Use Java *Double.longBitsToDouble* and *Double.doubleToLongBits* to convert before memory

COMPARISON INSTRUCTIONS (FCMPS, FCMPD)

- LEGv8 does not specify flag-setting conditions for IEEE-754 comparisons
- Use ARMv8's ones

IEEE-754 Relationship	ARM APSR Flags			
	N	Z	C	V
Equal	0	1	1	0
Less Than	1	0	0	0
Greater Than	0	0	1	0
Unordered (<i>At least one argument was NaN.</i>)	0	0	1	1

SHOWING THE REGISTERS

S0

0xc00a0000

Hex

Dec

S0

-2.15625

Hex

Dec

THE FINAL VIEW

1

MOVK

X0, #192

2

LSL

X0, X0, #8

3

MOVK

X0, #10

4

LSL

X0, X0, #16

5

SUBI

SP, SP, #16

6

STUR

X0, [SP, #0]

7

LDURS

S0, [SP, #0]

8

0x800000000:

0x0

Hex

0x7fffffff80:

0x0

Hex

0x7fffffff8:

0x0

Hex

0x7fffffff78:

0x0

Hex

0x7fffffff0:

0xc00a0000

Hex

0x7fffffff70:

0x0

Hex

0x7ffffffe8:

0x0

Hex

0x7ffffff68:

0x0

Hex

0x7ffffffe0:

0x0

Hex

0x7ffffff60:

0x0

Hex

0x7ffffffd8:

0x0

Hex

0x7ffffff58:

0x0

Hex

0x7ffffffd0:

0x0

Hex

0x7ffffff50:

0x0

Hex

0x7ffffffc8:

0x0

Hex

0x7ffffff48:

0x0

Hex

0x7ffffffc0:

0x0

Hex

0x7ffffff40:

0x0

Hex

0x7ffffffb8:

0x0

Hex

0x7ffffff38:

0x0

Hex

0x7ffffffb0:

0x0

Hex

0x7ffffff30:

0x0

Hex

0x7ffffffa8:

0x0

Hex

0x7ffffff28:

0x0

Hex

0x7ffffffa0:

0x0

Hex

0x7ffffff20:

0x0

Hex

0x7ffffff98:

0x0

Hex

0x7ffffff18:

0x0

Hex

0x7ffffff90:

0x0

Hex

0x7ffffff10:

0x0

Hex

0x7ffffff88:

0x0

Hex

0x7ffffff08:

0x0

Hex

Dec

Z

0

N

0

C

0

V

0

D0

0xc00a0000

Hex

Dec

D16

0x0

Hex

Dec

S0

-2.15625

Hex

Dec

S16

0x0

Dec

X16

0x0

Hex

Dec

D1

0x0

Hex

Dec

S1

0x0

Hex

Dec

S17

0x0

Dec

X17

0x0

Hex

Dec

D2

0x0

Hex

Dec

S2

0x0

Hex

Dec

S18

0x0

Dec

X18

0x0

Hex

Dec

D3

0x0

Hex

Dec

S3

0x0

Hex

Dec

S19

0x0

Dec

X19

0x0

Hex

Dec

D4

0x0

Hex

Dec

S4

0x0

Hex

Dec

S20

0x0

Dec

X20

0x0

Hex

Dec

D5

0x0

Hex

Dec

S5

0x0

Hex

Dec

S21

0x0

Dec

X21

0x0

Hex

Dec

D6

0x0

Hex

Dec

S6

0x0

Hex

Dec

S22

0x0

Dec

X22

0x0

Hex

Dec

D7

0x0

Hex

Dec

S7

0x0

Hex

Dec

S23

0x0

Dec

X23

0x0

Hex

Dec

D8

0x0

Hex

Dec

S8

0x0

Hex

Dec

S24

0x0

Dec

X24

0x0

Hex

Dec

D9

0x0

Hex

Dec

S9

0x0

Hex

Dec

S25

0x0

Dec

X25

0x0

Hex

Dec

D10

0x0

Hex

Dec

S10

0x0

Hex

Dec

S26

0x0

Dec

X26

0x0

Hex

Dec

D11

0x0

Hex

Dec

S11

0x0

Hex

Dec

S27

0x0

Dec

X27

0x0

Hex

Dec

D12

0x0

Hex

Dec

S12

0x0

Hex

Dec

S28

0x0

Dec

SP

0x7fffffff0

Hex

Dec

D13

0x0

Hex

Dec

S13

0x0

Hex

Dec

S29

0x0

Dec

FP

0x0

Hex

Dec

D14

0x0

Hex

Dec

S14

0x0

Hex

Dec

S30

0x0

Dec

LR

0x0

Hex

Dec

D15

0x0

Hex

Dec

S15

0x0

Hex

Dec

S31

0x0

Dec

XZR

0x0

Hex

Dec

Hex

Dec

Hex

Dec

THE CHERRY ON TOP: MODERNIZING THE BUILD SYSTEM

INTEGRATING MAVEN

- Latest GWT and AceGWT support Maven
- Integrated Maven into the simulator
- Can now use other IDEs, Java 21, GWT 2.11
- To develop, download the code and run *mvn package*. That's it.

CONCLUSIONS

- Arm's LEGv8 simulator finally working
- Only one to implement every LEGv8 instruction
- Can now be developed with modern tools, set-up and build in seconds

THANK YOU FOR
YOUR ATTENTION

[THESIS AVAILABLE HERE](#)

[SIMULATOR AVAILABLE HERE](#)