



UNIVERSITÀ  
DEGLI STUDI  
DI TRIESTE

# Restoration and development of Arm's Java-based LEGv8 ISA simulator

Graduating student - Simone Deiana

Supervisor - Alberto Carini

## Restoration and development of Arm's Java-based LEGv8 ISA simulator



## Restoration and development of Arm's Java-based LEGv8 ISA simulator





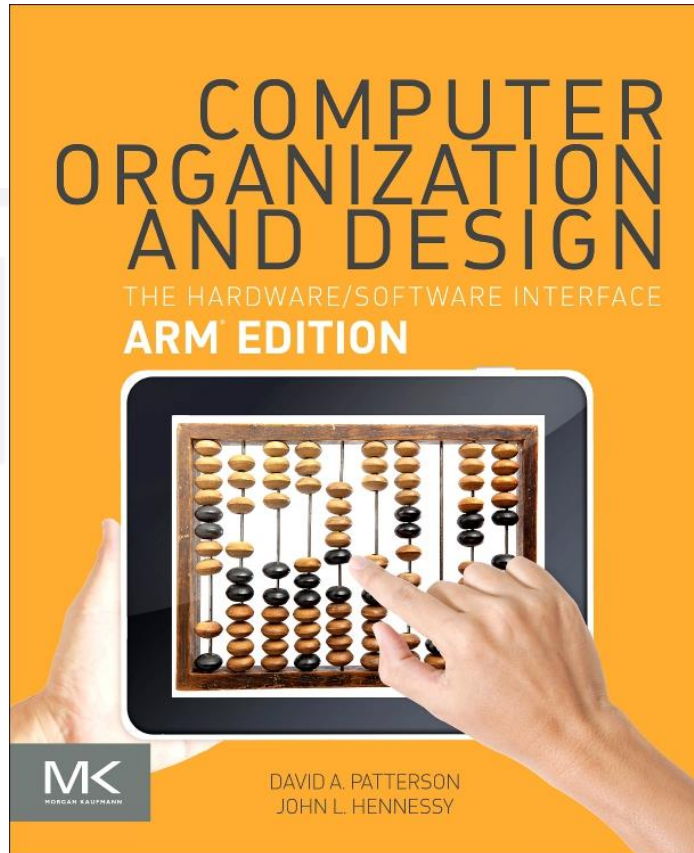
## Restoration and development of Arm's Java-based LEGv8 ISA simulator

# WHAT IS LEGv8?

Restoration and development of  
Arm's Java-based **LEGv8 ISA** simulator

WHAT IS LEGv8?

# AN ISA FOR LEARNING COMPUTER ARCHITECTURES



From Computer Organization and Design ARM Edition: The Hardware Software Interface - Patterson, D.A. and Hennessy, J.L.



David A.  
Patterson

Peg Skorpinski, CC BY-SA  
3.0  
<<https://creativecommons.org/licenses/by-sa/3.0/>>, via  
Wikimedia Commons



John L.  
Hennessy

Eric Chan, CC BY 2.0  
<<https://creativecommons.org/licenses/by/2.0/>>, via Wikimedia Commons



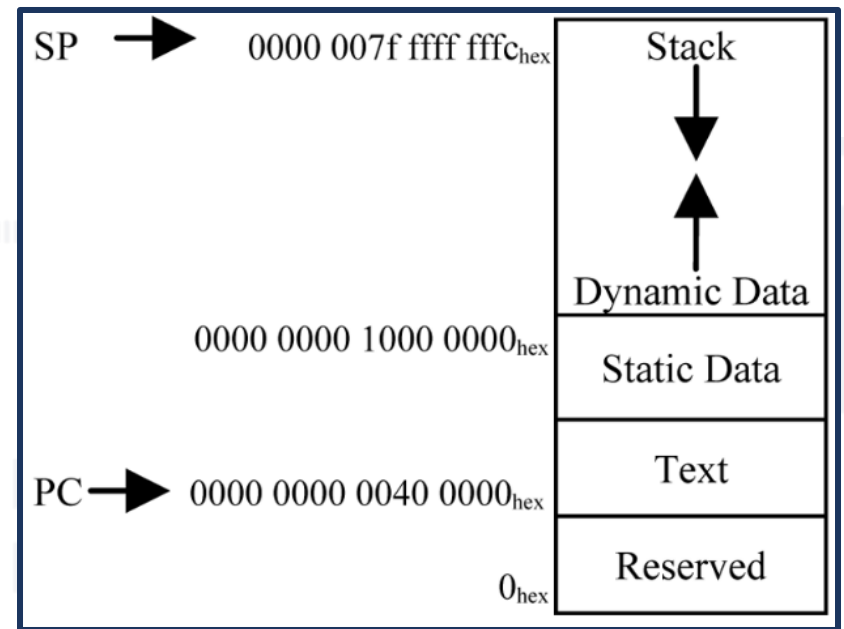
UNIVERSITÀ  
DEGLI STUDI  
DI TRIESTE

## THE DESIGN PHILOSOPHY

- As simple as it can be...
- ... but with a modern design
- Heavily inspired by ARMv8, almost a "subset"

# THE MEMORY

- 64-bit addresses
- Harvard model





## THE REGISTERS

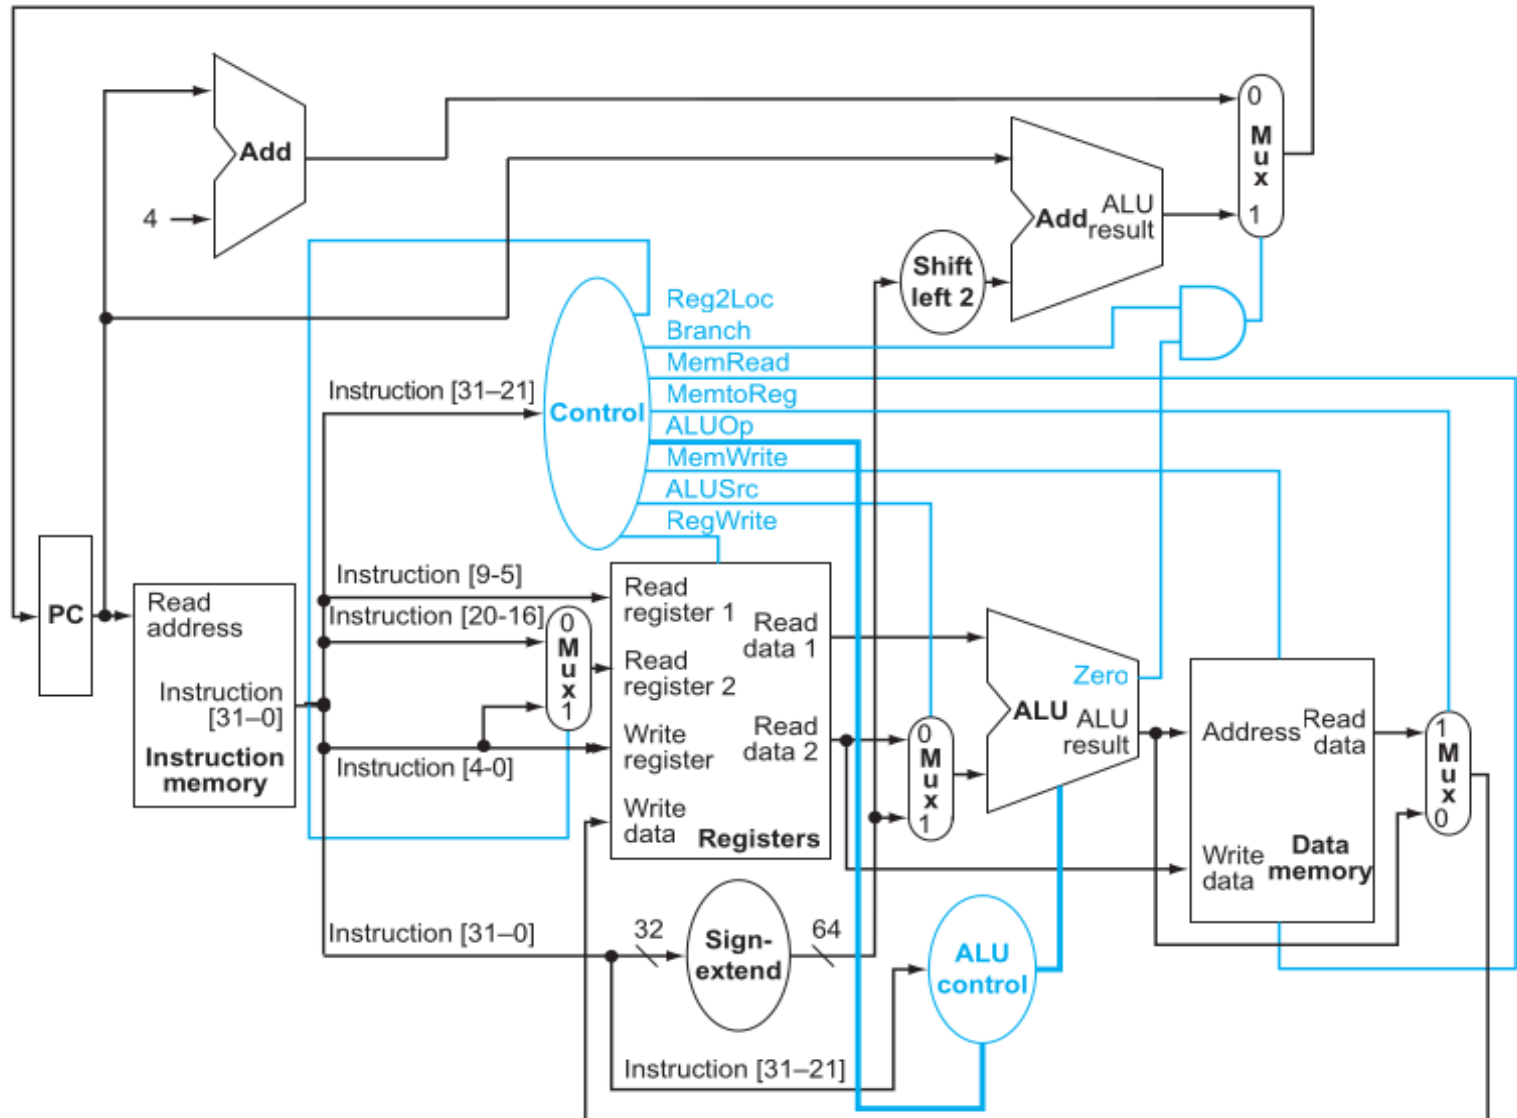
- 32 64-bit "X" integer registers
- 32 64-bit "D" floating-point registers
- 32 32-bit "S" floating-point "registers"

## THE INSTRUCTIONS

- 64-bit integer and IEEE-754 floating-point arithmetic
- Designed and optimized for pipelined execution

# WHAT IS LEGv8?

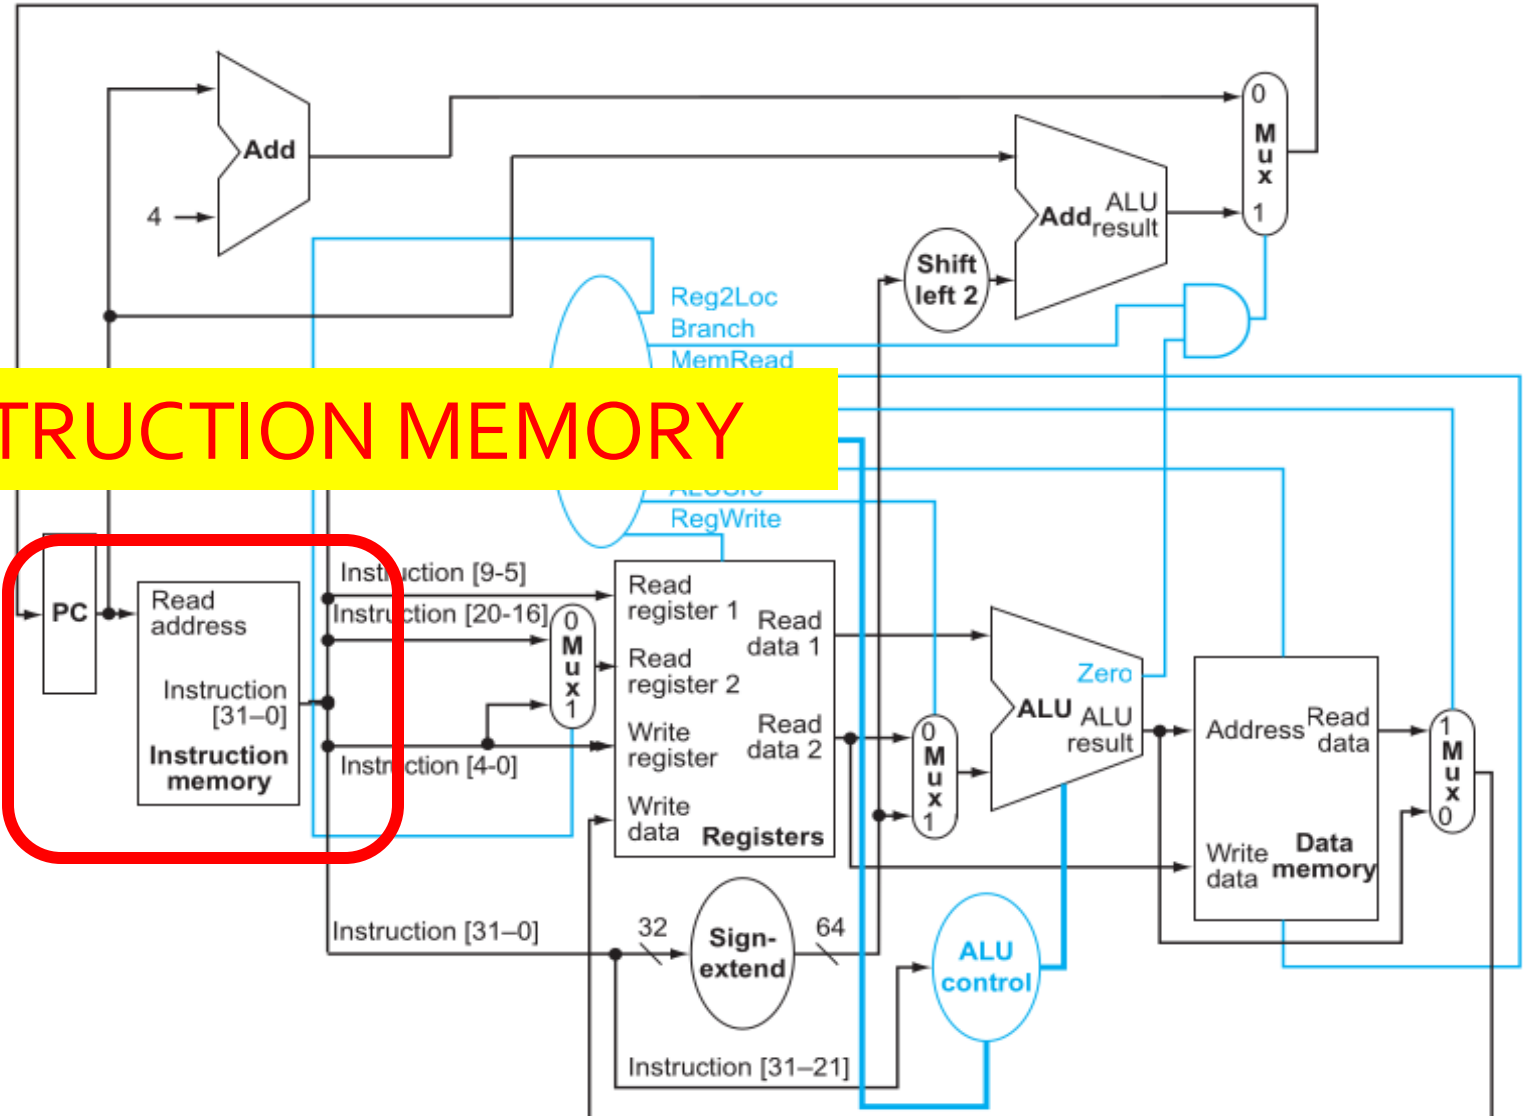
## THE DATAPATH



From Computer Organization and Design ARM Edition: The Hardware Software Interface - Patterson, D.A. and Hennessy, J.L.

# WHAT IS LEGv8?

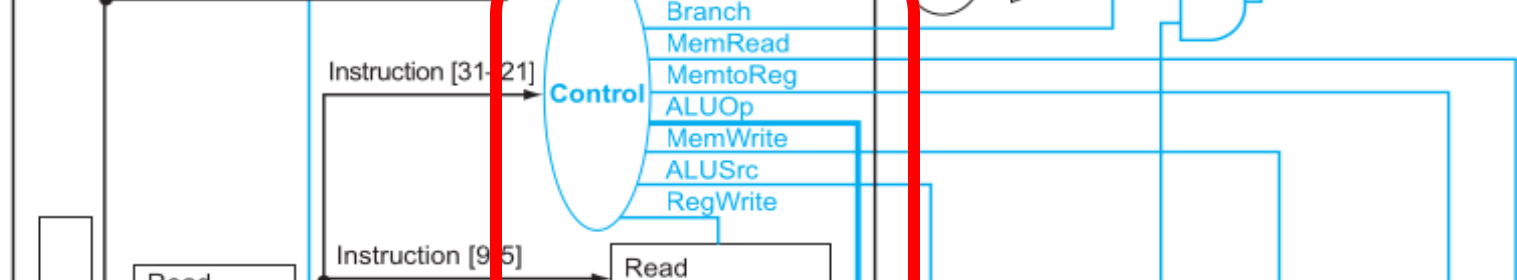
# THE DATAPAH



From Computer Organization  
and Design ARM Edition: The  
Hardware Software Interface -  
Patterson, D.A. and Hennessy,  
J.L.

Diagram of the Control Unit in a MIPS processor. The Control Unit is an oval labeled "Control". It receives two inputs: "Instruction [31-21]" and "Instruction [9-5]". It has seven outputs: "Branch", "MemRead", "MemtoReg", "ALUOp", "MemWrite", "ALUSrc", and "RegWrite". The "MemWrite" output is connected to a "Read" block.

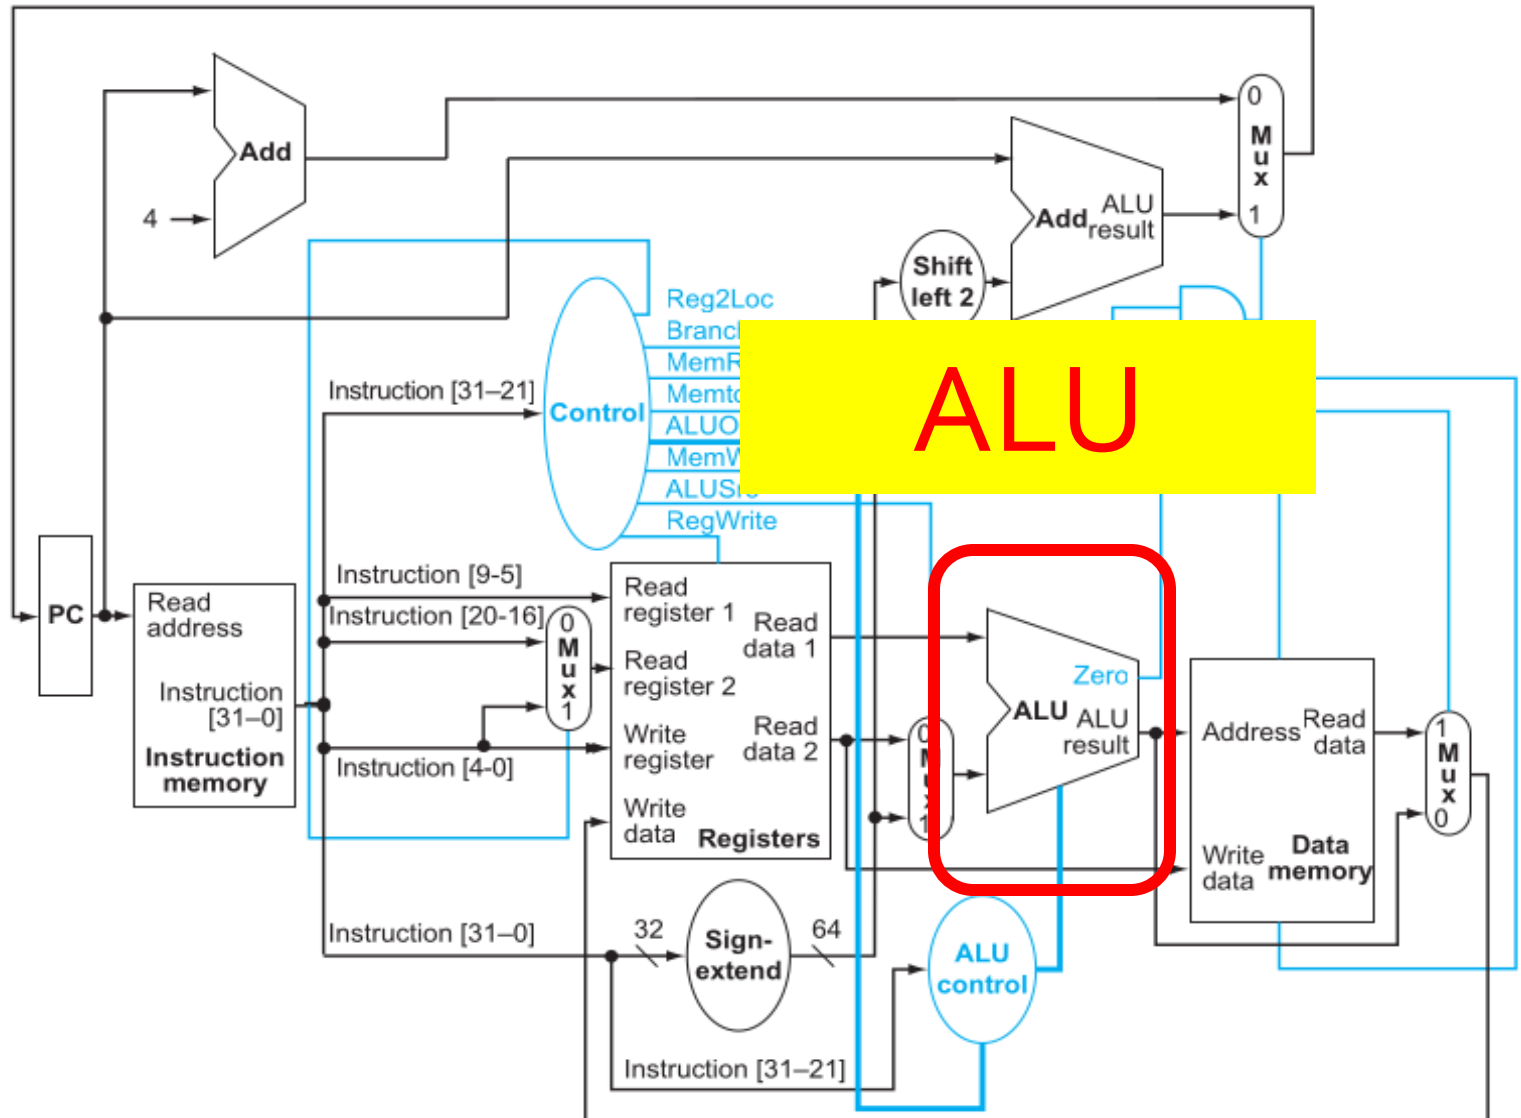
Diagram illustrating the Control Unit in a MIPS processor. The Control Unit is a central oval labeled "Control". It receives two inputs: "Instruction [31:21]" and "Instruction [9:5]". It outputs seven control signals: "Branch", "MemRead", "MemtoReg", "ALUOp", "MemWrite", "ALUSrc", and "RegWrite". The "RegWrite" signal is connected to a "Read" block, which also receives "Instruction [9:5]" as input.



From Computer Organization  
and Design ARM Edition: The  
Hardware Software Interface -  
Patterson, D.A. and Hennessy,  
J.L.

# WHAT IS LEGv8?

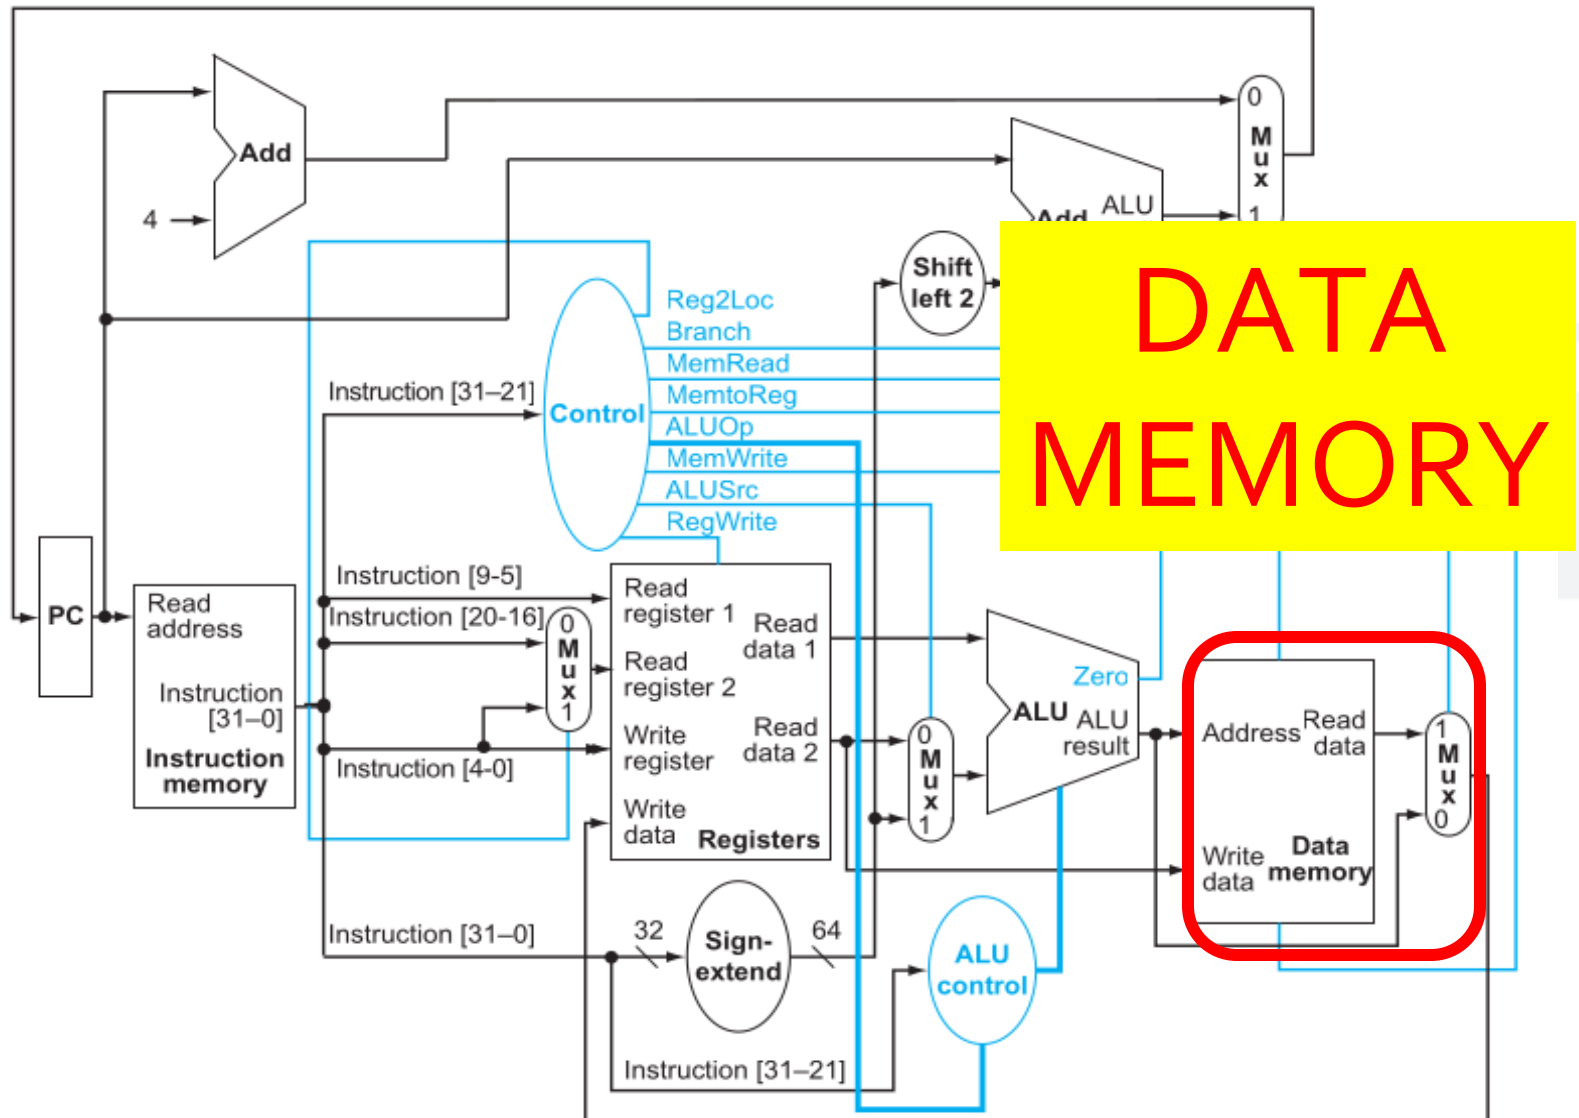
## THE DATAPATH



From Computer Organization and Design ARM Edition: The Hardware Software Interface - Patterson, D.A. and Hennessy, J.L.

# WHAT IS LEGv8?

## THE DATAPATH



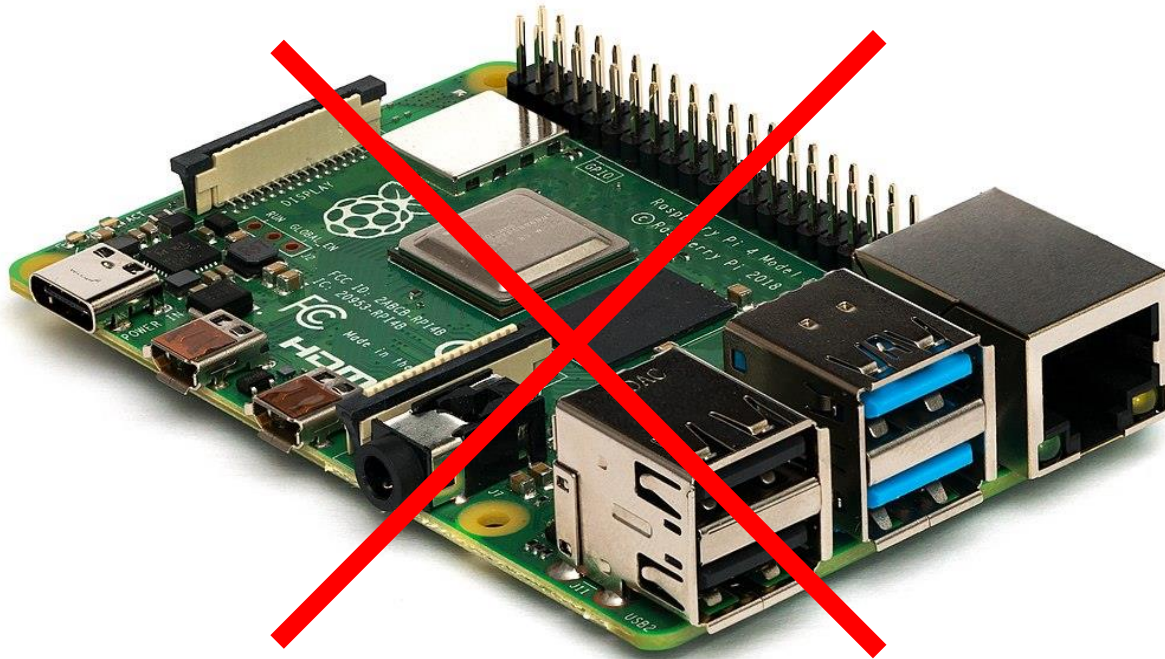
# WHAT SIMULATOR, AND WHY?

Restoration and development of  
Arm's Java-based LEGv8 ISA simulator



## WHAT SIMULATOR, AND WHY?

NO HARDWARE FOR LEGv8 => NEED A SIMULATOR



Michael H. („Laserlicht“) / Wikimedia Commons

WHAT SIMULATOR, AND WHY?

BUT WHICH ONE?



LEGv8

**197 results**

# THE PROBLEM:

NO SOFTWARE CAN YET  
SIMULATE THE ENTIRE  
LEGv8 ISA!

# THE SOLUTION:

- Write one from scratch
- OR (BETTER)
- Improve one that already exists

## THE GOOD NEWS

ARM HAS OFFICIALLY MADE A  
LEGv8 SIMULATOR

## WHAT STANDS OUT:

- Written in Java (high level, extensible)
- Distributed as a web application
- Nice, functional UI
- Closely follows the textbook

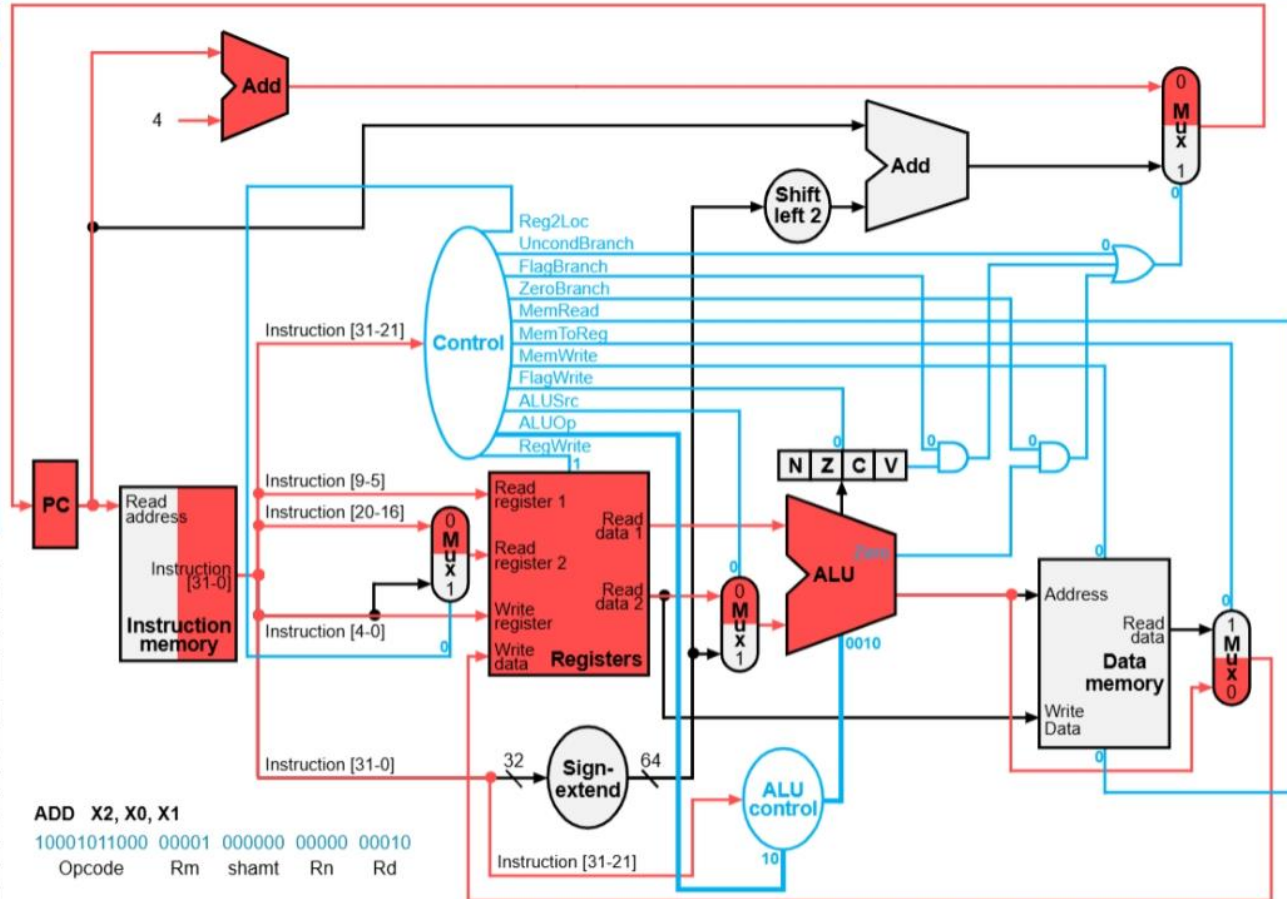
# WHAT SIMULATOR, AND WHY?

## LEGv8 Simulator

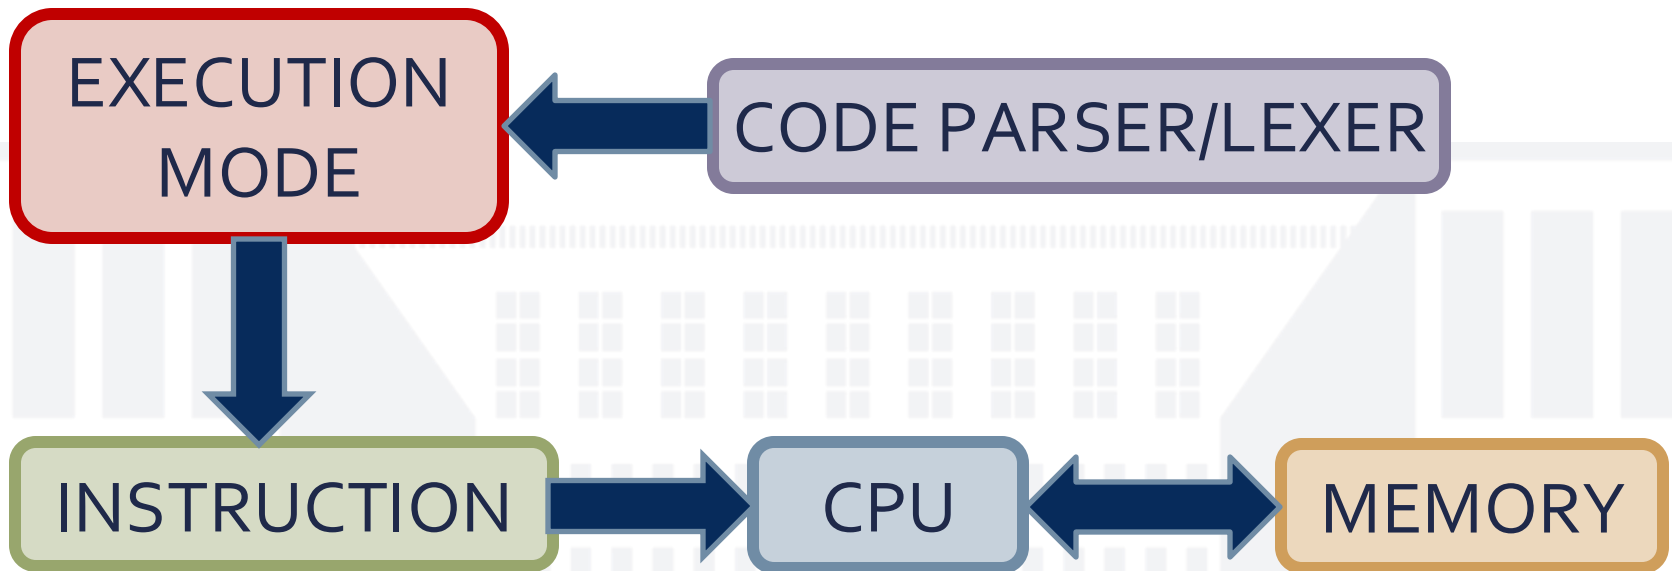
Execution Mode: **Single Cycle** Assemble Execute Instruction Help

```
1 ADDI X0, XZR, #53
2 ADDI X1, XZR, #75
3 ADD X2, X0, X1
4 CMP X2, X0
5 B.LE if
6 if: SUBIS X3, X2, #9
7 MOV X4, X3
8
```

PC	0x40000c	Hex	Dec	Z	0	N	0	C	0	V	0
X0	0x35	Hex	Dec	X16	0x0	Hex	Dec	X16	0x0	Hex	Dec
X1	0x4b	Hex	Dec	X17	0x0	Hex	Dec	X17	0x0	Hex	Dec
X2	0x80	Hex	Dec	X18	0x0	Hex	Dec	X18	0x0	Hex	Dec
X3	0x0	Hex	Dec	X19	0x0	Hex	Dec	X19	0x0	Hex	Dec
X4	0x0	Hex	Dec	X20	0x0	Hex	Dec	X20	0x0	Hex	Dec
X5	0x0	Hex	Dec	X21	0x0	Hex	Dec	X21	0x0	Hex	Dec
X6	0x0	Hex	Dec	X22	0x0	Hex	Dec	X22	0x0	Hex	Dec
X7	0x0	Hex	Dec	X23	0x0	Hex	Dec	X23	0x0	Hex	Dec
X8	0x0	Hex	Dec	X24	0x0	Hex	Dec	X24	0x0	Hex	Dec
X9	0x0	Hex	Dec	X25	0x0	Hex	Dec	X25	0x0	Hex	Dec
X10	0x0	Hex	Dec	X26	0x0	Hex	Dec	X26	0x0	Hex	Dec
X11	0x0	Hex	Dec	X27	0x0	Hex	Dec	X27	0x0	Hex	Dec
X12	0x0	Hex	Dec	SP	0x7ffffffc	Hex	Dec	SP	0x7ffffffc	Hex	Dec
X13	0x0	Hex	Dec	FP	0x0	Hex	Dec	FP	0x0	Hex	Dec
X14	0x0	Hex	Dec	LR	0x0	Hex	Dec	LR	0x0	Hex	Dec
X15	0x0	Hex	Dec	XZR	0x0	Hex	Dec	XZR	0x0	Hex	Dec



# STRUCTURE OF THE SIMULATOR





# THE BAD NEWS

IT'S INCOMPLETE  
AND BROKEN

# FIXING AND RESTORING THE SIMULATOR

Restoration and development of  
Arm's Java-based LEGv8 ISA simulator

# COMPARISONS DON'T WORK!

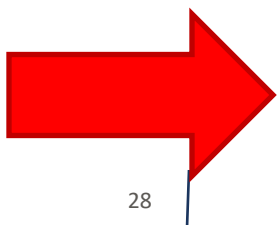
- No "if-else" conditionals
- No "switch-case" conditionals
- No "while" loops
- No "for" loops

## THE COMPARISON BUG

```
1      ADDI    X0, XZR, #1
2      ADDI    X1, XZR, #2
3      CMP     X0, X1
4      B.LE    should_jump
5      ADD     X2, X0, X1
6  should_jump:  ADDI    X2, XZR, #16
7
```

**X0 = 1, X1 = 2**

PC	0x400008	Hex	Dec	Z	0	N	0	C	0	V	0
X0	0x1	Hex	Dec	X16			0x0			Hex	Dec
X1	0x2	Hex	Dec	X17			0x0			Hex	Dec
X2	0x0	Hex	Dec	X18			0x0			Hex	Dec

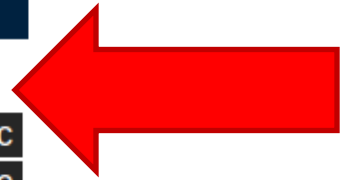


# THE COMPARISON BUG

1	ADDI	X0, XZR, #1
2	ADDI	X1, XZR, #2
3	CMP	X0, X1
4	B.LE	should_jump
5	ADD	X2, X0, X1
6	should_jump:	ADDI X2, XZR, #16
7		

COMPARE X0 WITH X1

PC	0x40000c	Hex	Dec	Z	0	N	1	C	0	V	1
X0	0x1	Hex	Dec	X16			0x0			Hex	Dec
X1	0x2	Hex	Dec	X17			0x0			Hex	Dec
X2	0x0	Hex	Dec	X18			0x0			Hex	Dec



## THE COMPARISON BUG

```
1 |      ADDI      X0, XZR, #1
2 |      ADDI      X1, XZR, #2
3 |      CMP       X0, X1
4 |      B.LE      should_jump
5 |      ADD       X2, X0, X1
6 |  should_jump:  ADDI      X2, XZR, #16
7 |
```

1 <= 2 ?

OF COURSE, LET'S JUMP!

PC	0x400010	Hex	Dec	Z	0	N	1	C	0	V	1
X0	0x1	Hex	Dec	X16			0x0	Hex	Dec		
X1	0x2	Hex	Dec	X17			0x0	Hex	Dec		
X2	0x0	Hex	Dec	X18			0x0	Hex	Dec		

## THE COMPARISON BUG

```
1 |      ADDI      X0, XZR, #1
2 |      ADDI      X1, XZR, #2
3 |      CMP       X0, X1
4 |      B.LE      should_jump
5 |      ADD       X2, X0, X1
6 | should_jump:   ADDI      X2, XZR, #16
7 |
```

...OR NOT

PC	0x400014	Hex	Dec	Z	0	N	1	C	0	V	1
X0	0x1	Hex	Dec	X16			0x0	Hex	Dec		
X1	0x2	Hex	Dec	X17			0x0	Hex	Dec		
X2	0x3	Hex	Dec	X18			0x0	Hex	Dec		

# BRANCH AND LINKS DON'T WORK!

~~*void subroutine(arg1, ...)*~~  
~~*float function(arg1, ...)*~~

CAN'T REUSE CODE



# THE BRANCH AND LINK BUG

```
1      ADDI      X0, XZR, #1
2      BL        subroutine
3      B         exit
4
5  subroutine:
6      ADDI      X0, X0, #16
7      BR        LR
8  exit:
9
```

# THE BRANCH AND LINK BUG

```
1      ADDI      X0, XZR, #1
2      BL        subroutine
3      B         exit
4
5  subroutine:
6      ADDI      X0, X0, #16
7      BR        LR
8  exit:
9
```

# THE BRANCH AND LINK BUG


```
1      ADDI    X0, XZR, #1
2      BL      subroutine
3      B       exit
4
5  subroutine:
6      ADDI    X0, X0, #16
7      BR      LR
8  exit:
9
```

# THE BRANCH AND LINK BUG

```
1      ADDI    X0, XZR, #1
2      BL      subroutine
3      B       exit
4
5  subroutine:
6      ADDI    X0, X0, #16
7      BR      LR
8  exit:
9
```

# THE BRANCH AND LINK BUG

1	ADDI	X0, XZR, #1	
2	BL	subroutine	
3	B	exit	
4			
5	subroutine:		
6	ADDI	X0, X0, #16	
7	BR	LR	
8	exit:		
9			

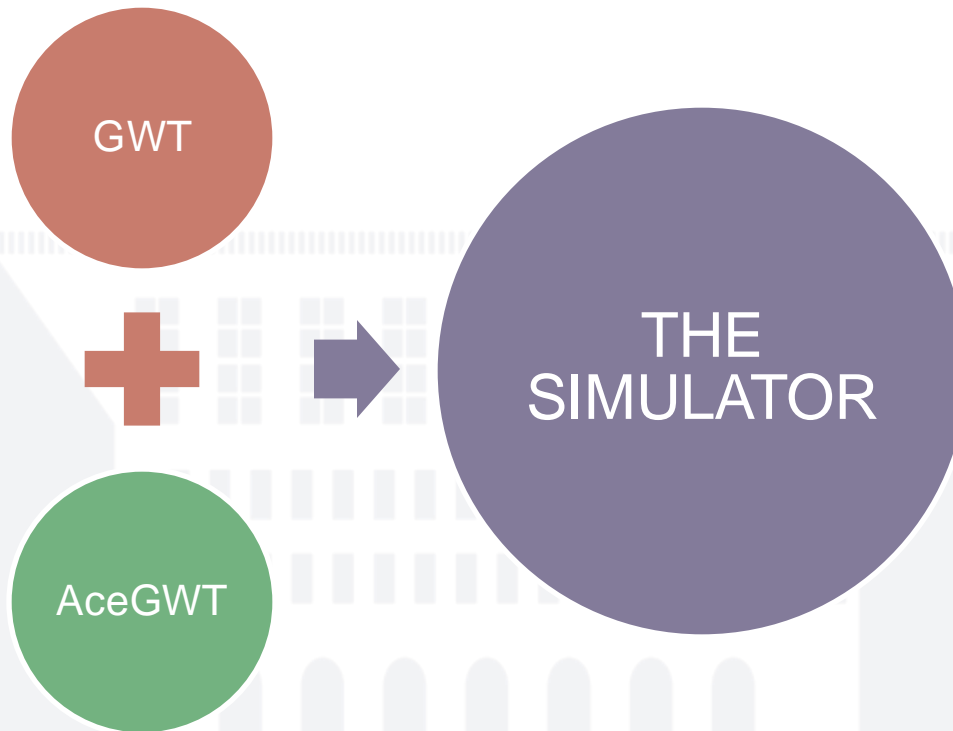


WE SHOULD BE HERE!

IT DOESN'T GO BACK!

ALL FIXED, BUT...  
NOBODY KNOWS HOW IT WORKS!

# THE PROJECT'S DEPENDENCIES



### GWT

- Framework (formerly) from Google
- Generates web applications (client-server, client only) from Java
- Emulates Java's JVM with JavaScript



# GWT

- Old, outdated, barely supported
- Convoluted custom build tools
- Limited emulation of JVM
- Basically needs Eclipse plug-in for real development

### AceGWT

- Provides GWT bindings for the Ace editor
- Can be used like normal GWT component
- Also old, outdated, and unsupported

### WORKING IT OUT:

- Needed older version of Eclipse, and Eclipse GWT plug-in
- Reverse engineered the dependencies and their configuration
- Configured build system to stop failing

AFTER 3 YEARS, WE CAN NOW  
PRODUCE NEW VERSIONS  
OF THE SIMULATOR!

# FILLING THE GAPS

Restoration and **development** of  
Arm's Java-based LEGv8 ISA simulator

# WHAT IS THE SIMULATOR MISSING?

- Incomplete integer arithmetic
- No visualization for the stack memory
- No IEEE-754 arithmetic and data instructions

# THE MISSING INTEGER-BASED INSTRUCTIONS

- **MUL** — LOWER 64 BITS OF THE MULTIPLICATION
- **SMULH** — HIGHER 64 BITS OF THE SIGNED MULTIPLICATION
- **UMULH** — HIGHER 64 BITS OF THE UNSIGNED MULTIPLICATION
- **SDIV** — SIGNED DIVISION
- **UDIV** — UNSIGNED DIVISION
- **LDA** — LOAD ADDRESS OF A LABEL IN A REGISTER

Easy to implement in  
the existing codebase

but...

Java doesn't like big  
or unsigned numbers!



## PROBLEM 1

- Java does not have primitive 128-bit integer types
- The BigInteger library exists
- GWT 2.7 doesn't emulate it (2.8 does)

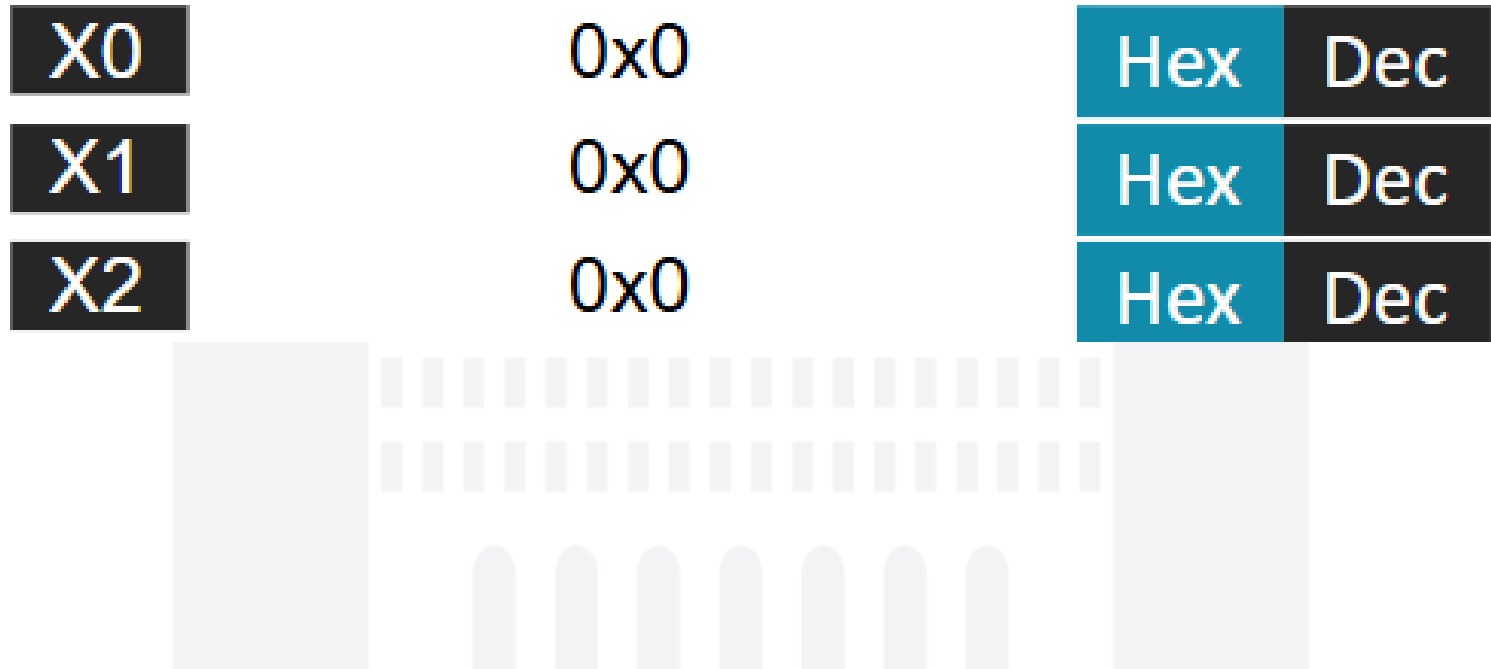
## PROBLEM 2

- Primitive integers are signed
- BigInteger also signed
- Bitmask converts 64-bit unsigned integers to 65-bit signed, perform signed operations, back to unsigned

# STACK NOT VISUALIZED

- Important for testing and debugging complex programs (now we can write them)
- Useful to understand LEGv8 and stack management

# TAKING SOME INSPIRATION



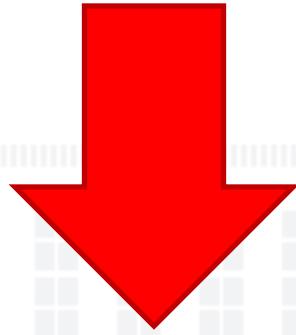
## THE NEW STACK VIEW

0x8000000000:	0x0	Hex	0x7fffffff80:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff78:	0x0	Hex
0x7fffffff0:	0x0	Hex	0x7fffffff70:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff68:	0x0	Hex
0x7fffffff80:	0x0	Hex	0x7fffffff60:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff58:	0x0	Hex
0x7fffffff80:	0x0	Hex	0x7fffffff50:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff48:	0x0	Hex
0x7fffffff80:	0x0	Hex	0x7fffffff40:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff38:	0x0	Hex
0x7fffffff80:	0x0	Hex	0x7fffffff30:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff28:	0x0	Hex
0x7fffffff80:	0x0	Hex	0x7fffffff20:	0x0	Hex
0x7fffffff8:	0x0	Hex	0x7fffffff18:	0x0	Hex
0x7fffffff80:	0x0	Hex	0x7fffffff10:	0x0	Hex
0x7fffffff88:	0x0	Hex	0x7fffffff08:	0x0	Hex

# ADDING FLOATING-POINT SUPPORT

- **FADDS, FADDD** — ADD TWO IEEE-754 VALUES
- **FSUBS, FSUBD** - SUBTRACT TWO IEEE-754 VALUES
- **FMULS, FMULD** - MULTIPLY TWO IEEE-754 VALUES
- **FDIVS, FDIVD** — DIVIDE TWO IEEE-754 VALUES
- **LDURS, LDURD** — LOAD IEEE-754 VALUE FROM MEMORY
- **STURS, STURD** - STORE IEEE-754 VALUE TO MEMORY
- **FCMPS, FCMPPD** — COMPARE TWO IEEE-754 VALUES

# SIMULATOR NOT DESIGNED FOR FLOATING POINT



SOME STRUCTURAL CHANGES NEEDED  
(e.g. parser)

# ARITHMETICAL INSTRUCTIONS





# MEMORY ACCESS INSTRUCTIONS

- Simulator designed for integer use
- Memory uses integers to store bytes

# IDEA: SEE INTEGERS LIKE RAW BITS



CONVERT FLOATS  
TO RAW BITS

CAN BE USED AS  
BINARY PROTOCOL



RETROCOMPATIBILITY!

# COMPARISON INSTRUCTIONS

- LEGv8 does not specify flag-setting conditions for IEEE-754 comparisons
- Use ARMv8's ones:

IEEE-754 Relationship	ARM APSR Flags			
	N	Z	C	V
Equal	0	1	1	0
Less Than	1	0	0	0
Greater Than	0	0	1	0
Unordered ( <i>At least one argument was NaN.</i> )	0	0	1	1

## THE FINAL VIEW

1	MOVK	X0, #192	0x800000000:	0x0	Hex	0x7fffffff80:	0x0	Hex
2	LSL	X0, X0, #8	0x7fffffff8:	0x0	Hex	0x7fffffff78:	0x0	Hex
3	MOVK	X0, #10	0x7fffffff0:	0xc00a0000	Hex	0x7fffffff70:	0x0	Hex
4	LSL	X0, X0, #16	0x7fffffff8:	0x0	Hex	0x7fffffff68:	0x0	Hex
5	SUBI	SP, SP, #16	0x7ffffffe0:	0x0	Hex	0x7fffffff60:	0x0	Hex
6	STUR	X0, [SP, #0]	0x7ffffffd8:	0x0	Hex	0x7fffffff58:	0x0	Hex
7	LDURS	S0, [SP, #0]	0x7ffffffd0:	0x0	Hex	0x7fffffff50:	0x0	Hex
8			0x7ffffffc8:	0x0	Hex	0x7fffffff48:	0x0	Hex
			0x7ffffffc0:	0x0	Hex	0x7fffffff40:	0x0	Hex
			0x7ffffffb8:	0x0	Hex	0x7fffffff38:	0x0	Hex
			0x7ffffffb0:	0x0	Hex	0x7fffffff30:	0x0	Hex
			0x7ffffffa8:	0x0	Hex	0x7fffffff28:	0x0	Hex
			0x7ffffffa0:	0x0	Hex	0x7fffffff20:	0x0	Hex
			0x7fffffff98:	0x0	Hex	0x7fffffff18:	0x0	Hex
			0x7fffffff90:	0x0	Hex	0x7fffffff10:	0x0	Hex
			0x7fffffff88:	0x0	Hex	0x7fffffff08:	0x0	Hex

PC	0x40001c	Z	0	N	0	C	0	V	0	D0	0xc00a0000	D16	0x0	S0	-2.15625	S16	0x0
X0	0xc00a0000	Hex	Dec	X16	0x0	Hex	Dec	D1	0x0	Hex	Dec	D17	0x0	Hex	Dec	S17	0x0
X1	0x0	Hex	Dec	X17	0x0	Hex	Dec	D2	0x0	Hex	Dec	D18	0x0	Hex	Dec	S18	0x0
X2	0x0	Hex	Dec	X18	0x0	Hex	Dec	D3	0x0	Hex	Dec	D19	0x0	Hex	Dec	S19	0x0
X3	0x0	Hex	Dec	X19	0x0	Hex	Dec	D4	0x0	Hex	Dec	D20	0x0	Hex	Dec	S20	0x0
X4	0x0	Hex	Dec	X20	0x0	Hex	Dec	D5	0x0	Hex	Dec	D21	0x0	Hex	Dec	S21	0x0
X5	0x0	Hex	Dec	X21	0x0	Hex	Dec	D6	0x0	Hex	Dec	D22	0x0	Hex	Dec	S22	0x0
X6	0x0	Hex	Dec	X22	0x0	Hex	Dec	D7	0x0	Hex	Dec	D23	0x0	Hex	Dec	S23	0x0
X7	0x0	Hex	Dec	X23	0x0	Hex	Dec	D8	0x0	Hex	Dec	D24	0x0	Hex	Dec	S24	0x0
X8	0x0	Hex	Dec	X24	0x0	Hex	Dec	D9	0x0	Hex	Dec	D25	0x0	Hex	Dec	S25	0x0
X9	0x0	Hex	Dec	X25	0x0	Hex	Dec	D10	0x0	Hex	Dec	D26	0x0	Hex	Dec	S26	0x0
X10	0x0	Hex	Dec	X26	0x0	Hex	Dec	D11	0x0	Hex	Dec	D27	0x0	Hex	Dec	S27	0x0
X11	0x0	Hex	Dec	X27	0x0	Hex	Dec	D12	0x0	Hex	Dec	D28	0x0	Hex	Dec	S28	0x0
X12	0x0	Hex	Dec	SP	0x7fffffff0	Hex	Dec	D13	0x0	Hex	Dec	D29	0x0	Hex	Dec	S29	0x0
X13	0x0	Hex	Dec	FP	0x0	Hex	Dec	D14	0x0	Hex	Dec	D30	0x0	Hex	Dec	S30	0x0
X14	0x0	Hex	Dec	LR	0x0	Hex	Dec	D15	0x0	Hex	Dec	D31	0x0	Hex	Dec	S31	0x0
X15	0x0	Hex	Dec	XZR	0x0	Hex	Dec										

# THE CHERRY ON TOP: MODERNIZING THE BUILD SYSTEM

## WHY DO THIS?

- Latest GWT and AceGWT support Maven
- Much easier to include if project also supports Maven

## THE TRICKLE-DOWN EFFECT

- Libraries are managed automatically
- Project decoupled from Eclipse, can use other IDEs or even terminal
- Can use Java 21, GWT 2.11
- More configurable builds

# CONCLUSIONS

- Arm's LEGv8 simulator finally working
- Only one to implement every LEGv8 instruction
- Can now be developed with modern tools, set-up and build in seconds, much easier collaboration



THANK YOU FOR  
YOUR ATTENTION

[THESIS AVAILABLE HERE](#)

[SIMULATOR AVAILABLE HERE](#)